

Starting and Stopping Services

IN THIS CHAPTER

Understanding the various Linux init daemons

Auditing Linux init daemon-controlled services

Stopping and starting services

Changing the Linux server's default runlevel

Removing services

A *service*, also called a daemon, is a running program or a process that provides a particular function. Your Linux server's primary job is to offer services. As a Linux server administrator, one of your primary jobs is managing these services.

The `cupsd` daemon is an example of a service your Linux server can provide. This service is responsible for printing and is covered in Chapter 16, “Configuring a Print Server.”

How do you start the `cupsd` daemon? How do you stop it? How do you get it to be started at system boot time? What if you have a special program you want to start on your server? How do you get it to start at boot time? All these questions and more will be answered in this chapter. Soon, you will have an excellent grasp on managing these services.

Understanding the Linux init Daemon

In order to understand service management, you need to understand the `init` daemon. The `init` daemon can be thought of as the “mother of all processes.” This daemon is the first process to be started by the kernel on the Linux server. The Linux kernel has a process ID (PID) of 0. Thus, the `init` daemon has a parent process ID (PPID) of 0, and a PID of 1. Once started, `init` is responsible for spawning (launching) processes configured to be started at the server's boot time, such as the login shell (`getty` or `mingetty` process). It is also responsible for managing services. That's a lot of responsibility for one daemon!

The Linux `init` daemon was based upon the Unix System V `init` daemon. Thus, it is called the `SysVinit` daemon. However, it was not the only classic `init` daemon. The `init` daemon is not

part of the Linux kernel. Therefore, it can come in different flavors, and Linux distributions can choose which flavor to use. Another classic `init` daemon was based on Berkeley Unix, also called BSD. Therefore, the two original Linux `init` daemons were BSD `init` and `SysVinit`.

The classic `init` daemons worked without problems for many years. However, these daemons were created to work within a static environment. As new hardware, such as USB devices, came along, the classic `init` daemons had trouble dealing with these and other hot-plug devices. Computer hardware had changed from static to event-based. New `init` daemons were needed to deal with these fluid environments.

In addition, as new services came along, the classic `init` daemons had to deal with starting more and more services. Thus, the entire system initialization process was less efficient and ultimately slower.

The modern `init` daemons have attempted to solve the problems of inefficient system boots and non-static environments. Two of these `init` daemons are `Upstart` and `systemd`. Many Linux distributions have made the move to the newer `init` daemons while maintaining backward compatibility to the classic `SysVinit` or BSD `init` daemons.

`Upstart`, available at <http://upstart.ubuntu.com>, was originally developed by Canonical, the parent of the Ubuntu distribution. However, many other distributions have adopted it, including:

- RHEL version 6
- Fedora versions 9 through 14
- Ubuntu versions 6–10 and up
- openSUSE versions 11.3 and up

A new daemon, `systemd`, available at <http://fedoraproject.org/wiki/Systemd>, was written primarily by Lennart Poettering, a Red Hat developer. It is currently used by Fedora versions 15 and up.

In order to properly manage your services, you need to know which `init` daemon your server has. Figuring that out can be a little tricky. Look through the following to help determine your Linux server's `init` daemon.

- Do your Linux distribution and version appear in the preceding list of `Upstart` adopters? Then your Linux `init` daemon is the `Upstart` `init` daemon.
- Is your Linux distribution Fedora version 15 or above? Then your Linux `init` daemon is the `systemd` `init` daemon.
- Try searching your Linux distribution's `init` daemon for clues, using the `strings` and the `grep` commands. The following code example shows the `init` daemon on a Linux Mint distribution being searched for `systemd` and `Upstart` `init` daemon references. The search for `systemd` yields nothing. However, the

search for Upstart produces results. Thus, you can see that this Linux Mint distribution uses the Upstart `init` daemon.

```
$ sudo strings /sbin/init | grep -i systemd
$
$ sudo strings /sbin/init | grep -i upstart
upstart-devel@lists.ubuntu.com
UPSTART_CONFDIR
UPSTART_NO_SESSIONS
...
```

On a Fedora server, the search for Upstart yields nothing. However, you can see the search for `systemd` yields the existence of the `systemd` daemon.

```
# strings /sbin/init | grep -i upstart
#
# strings /sbin/init | grep -i systemd
systemd.unit=
systemd.log_target=
systemd.log_level=
...
```

Tip

If you do not have the `strings` command on your Linux system, you can install it via the `binutils` package. On RHEL and Fedora, use the command `yum install binutils`. On Ubuntu, use the command `sudo apt-get install binutils`.

- If you still cannot tell what `init` daemon your server has, try looking on the `init` Wikipedia page (<http://wikipedia.org/wiki/Init>) under “Replacements for `init`?”
- Have you been unable to find any information using the preceding suggestions? Then, most likely, your distribution is still using the classic `SysVinit` or `BSD init` daemon.

Keep in mind that some Linux distributions have not moved to the newer daemons. Most of those that have, do maintain backward compatibility with the `SysVinit` and `BSD init` daemons.

Understanding the classic `init` daemons

The classic `init` daemons, `SysVinit` and `BSD init`, are worth understanding, even if your Linux server has a different `init` daemon. Not only is backward compatibility to the classics often used in the newer `init` daemons, but many are based upon them. Understanding the classic `init` daemons will help you to understand the modern `init` daemons.

The classic `SysVinit` and `BSD init` daemons operate in a very similar fashion. Although in the beginning they may have been rather different, over time, very few

significant differences remained. For example, the older BSD `init` daemon would obtain configuration information from the `/etc/ttytab` file. Now, like the `SysVinit` daemon, the BSD `init` daemon's configuration information is taken at boot time from the `/etc/inittab` file. The following is a classic `SysVinit` `/etc/inittab` file:

```
# cat /etc/inittab
#
# inittab          This file describes how the INIT process should set up
#                  the system in a certain run-level.
#
# Author:          Miquel van Smoorenburg, <miquels@drinkel.nl.mugnet.org>
#                  Modified for RHS Linux by Marc Ewing and Donnie Barnes
#

# Default runlevel. The runlevels used by RHS are:
#  0 - halt (Do NOT set initdefault to this)
#  1 - Single user mode
#  2 - Multiuser, without NFS (The same as 3, if you do not have networking)
#  3 - Full multiuser mode
#  4 - unused
#  5 - X11
#  6 - reboot (Do NOT set initdefault to this)
#
id:5:initdefault:

# System initialization.
si::sysinit:/etc/rc.d/rc.sysinit

10:0:wait:/etc/rc.d/rc 0
11:1:wait:/etc/rc.d/rc 1
12:2:wait:/etc/rc.d/rc 2
13:3:wait:/etc/rc.d/rc 3
14:4:wait:/etc/rc.d/rc 4
15:5:wait:/etc/rc.d/rc 5
16:6:wait:/etc/rc.d/rc 6

# Trap CTRL-ALT-DELETE
ca::ctrlaltdel:/sbin/shutdown -t3 -r now

# When our UPS tells us power has failed, assume we have a few minutes
# of power left.  Schedule a shutdown for 2 minutes from now.
# This does, of course, assume you have powerd installed and your
# UPS connected and working correctly.
pf::powerfail:/sbin/shutdown -f -h +2 "Power Failure; System Shutting Down"

# If power was restored before the shutdown kicked in, cancel it.
pr:12345:powerokwait:/sbin/shutdown -c "Power Restored; Shutdown Cancelled"
```

```
# Run gettys in standard runlevels
1:2345:respawn:/sbin/mingetty tty1
2:2345:respawn:/sbin/mingetty tty2
3:2345:respawn:/sbin/mingetty tty3
4:2345:respawn:/sbin/mingetty tty4
5:2345:respawn:/sbin/mingetty tty5
6:2345:respawn:/sbin/mingetty tty6

# Run xdm in runlevel 5
x:5:respawn:/etc/X11/prefdm -nodaemon
```

The `/etc/inittab` file tells the `init` daemon which runlevel is the default runlevel. A runlevel is a categorization number that determines what services are started and what services are stopped. In the preceding example, a default runlevel of 5 is set with the code `id:5:initdefault:`. Table 15.1 shows the standard seven Linux runlevels.

TABLE 15.1 Standard Linux Runlevels

Runlevel #	Name	Description
0	Halt	All services are shut down and the server is stopped.
1 or S	Single User Mode	The root account is automatically logged in to the server. Other users cannot log in to the server. Only the command line interface is available. Network services are not started.
2	Multuser Mode	Users can log in to the server, but only the command line interface is available. Network services are not started.
3	Extended Multuser Mode	Users can log in to the server, but only the command line interface is available. Network services are started.
4	User Defined	Users can customize this runlevel.
5	Graphical Mode	Users can log in to the server. Command line and graphical interfaces are available. Network services are started.
6	Reboot	The server is rebooted.

Linux distributions can differ slightly on the definition of each runlevel as well as which runlevels are offered. The Ubuntu distribution, for example, offers runlevels 0–6, but runlevels 2–5 start the same services as standard runlevel 5 listed in Table 15.1.

CAUTION

The only runlevels that should be used in the `/etc/inittab` file are 2 through 5. The other runlevels could cause problems. For example, if you put runlevel 6 in the `/etc/inittab` file as the default, when the server reboots, it would go into a loop and continue to reboot over and over again.

Notice that it doesn't make sense to include some of the `runlevel` categories described in Table 15.1 in the `/etc/inittab` file. There is no reason why you would want your default `runlevel` to be 0, which would halt the server every time it had just started.

The `runlevels` are not only used as a default `runlevel` in the `/etc/inittab` file. They can also be called directly using the `init` daemon itself. Thus, if you want to immediately halt your server, you type `init 0` at the command line:

```
# init 0
...
System going down for system halt NOW!
```

The `init` command will accept any of the `runlevel` numbers in Table 15.1, allowing you to quickly switch your server from one `runlevel` category to another. For example, if you need to perform troubleshooting, which requires the graphical interface to be down, you can type `init 3` at the command line:

```
# init 3
INIT: Sending processes the TERM signal
starting irqbalance:                [ OK ]
Starting setroubleshootd:
Starting fuse: Fuse filesystem already available.
...
Starting console mouse services:    [ OK ]
```

To see your Linux server's current `runlevel`, simply type in the command `runlevel`. The first item displayed will be the server's previous `runlevel`, which in the following example is 5. The second item displayed shows the server's current `runlevel`, which in this example is 3.

```
$ runlevel
5 3
```

In addition to the `init` command, you can also use the `telinit` command, which is functionally the same. In the example that follows, the `telinit` command is used to reboot the server by taking it to `runlevel 6`:

```
# telinit 6
INIT: Sending processes the TERM signal
Shutting down smartd:                [ OK ]
Shutting down Avahi daemon:          [ OK ]
Stopping dhcdd:                      [ OK ]
Stopping HAL daemon:                 [ OK ]
...
Starting killall:
Sending all processes the TERM signal... [ OK ]
Sending all processes the KILL signal... [ OK ]
```

```
...
Unmounting filesystems                [ OK ]
Please stand by while rebooting the system
...
```

On a freshly booted Linux server, the current runlevel number should be the same as the default runlevel number in the `/etc/inittab` file. However, notice that the previous runlevel in the example that follows is N. The N stands for “Nonexistent” and indicates the server has been freshly booted.

```
$ runlevel
N 5
```

How does the server know which services to stop and which ones to start when a particular runlevel is chosen? When a runlevel is chosen, the scripts located in the `/etc/rc.d/rc#.d` directory (where # is the chosen runlevel) are run. These scripts are run whether the runlevel is chosen via a server boot and the `/etc/inittab` runlevel setting, or when the `init` or `telinit` command is used. For example, if `runlevel 5` is chosen, then all the scripts in the `/etc/rc.d/rc5.d` directory are run.

```
# ls /etc/rc.d/rc5.d
K01smolt                K88wpa_supplicant    S22messagebus
K02avahi-dnscfg         K89dund              S25bluetooth
K02NetworkManager      K89netplugd         S25fuse
K02NetworkManagerDispatcher K89pand             S25netfs
K05ssaslauthd          K89rdisc            S25pcscd
K10dc_server           K91capi             S26hidd
K10psacct              S00microcode_ctl    S26udev-post
K12dc_client           S04readahead_early  S28autofs
K15gpm                 S05kudzu            S50hplip
K15httpd               S06cpuspeed         S55cups
K20nfs                 S08ip6tables        S55sshd
K24irda                S08iptables         S80sendmail
K25squid               S09isdn             S90ConsoleKit
K30spamassassin        S10network          S90crond
K35vncserver           S11auditd           S90xfs
K50netconsole          S12restorecond      S95anacron
K50tux                 S12syslog           S95atd
K69rpcsvcgssd          S13irqbalance       S96readahead_later
K73winbind             S13mcstrans         S97dhcdbd
K73ypbind              S13rpcbind          S97yum-updatesd
K74nsd                 S13setroubleshoot  S98avahi-daemon
K74ntpd                S14nfslock          S98haldaemon
K84btseed              S15mdmonitor        S99firstboot
K84bttrack             S18rpcidmapd        S99local
K87multipathd          S19rpcgssd          S99smartd
```

Notice that some of the scripts within the `/etc/rc.d/rc5.d` directory start with a `K` and some start with an `S`. The `K` refers to a script that will kill (stop) a process. The `S` refers to a script that will start a process. Also, each `K` and `S` script has a number before the name of the service or daemon they control. This allows the services to be stopped or started in a particular controlled order. You would not want your Linux server's network services to be started before the network itself was started.

An `/etc/rc.d/rc#.d` directory exists for all of the standard Linux runlevels. Each one contains scripts to start and stop services for its particular runlevel.

```
# ls -d /etc/rc.d/rc?.d
/etc/rc.d/rc0.d /etc/rc.d/rc2.d /etc/rc.d/rc4.d /etc/rc.d/rc6.d
/etc/rc.d/rc1.d /etc/rc.d/rc3.d /etc/rc.d/rc5.d
```

Actually, the files in the `/etc/rc.d/rc#.d` directories are not scripts, but instead symbolic links to scripts in the `/etc/rc.d/init.d` directory. Thus, there is no need to have multiple copies of particular scripts.

```
# ls -l /etc/rc.d/rc5.d/K15httpd
lrwxrwxrwx 1 root root 15 2011-10-27 19:59
/etc/rc.d/rc5.d/K15httpd -> ../init.d/httpd
#
# ls /etc/rc.d/init.d
anacron          functions        multipathd       rpcidmapd
atd              fuse            netconsole       rpcsvcgssd
auditd           gpm             netfs            saslauthd
autofs           haldaemon       netplugd         sendmail
avahi-daemon     halt            network          setroubleshoot
avahi-dnssconfd hidd            NetworkManager  single
bluetooth        hplip           NetworkManagerDispatcher smartd
btseed           hsqldb          nfs              smolt
bttrack          httpd           nfslock          spamassassin
capi             ip6tables       nscd             squid
ConsoleKit       iptables        ntpd             sshd
cpuspeed         irda            pand             syslog
crond            irqbalance      pcscd            tux
cups             isdn            psacct           udev-post
cups-config-daemon killall          rdisc            vncserver
dc_client        kudzu           readahead_early  winbind
dc_server        mcstrans        readahead_later  wpa_supplicant
dhcdbd          mdmonitor       restorecond      xfs
dund             messagebus      rpcbind          ypbind
firstboot        microcode_ctl   rpcgssd          yum-updatesd
```

Notice that each service has a single script in `/etc/rc.d/init.d`. There aren't separate scripts for stopping and starting a service. These scripts will stop or start a service depending upon what parameter is passed to them by the `init` daemon.

NOTE

Depending upon your distribution and your `init` daemon, a service's stop and start scripts may be stored in a different location than the `/etc/rc.d/init.d` directory. The three possible locations are:

- `/etc/rc.d/init.d` – A directory used by the SysVinit daemon.
- `/etc/rc.d` – The directory used by the BSD `init` daemon
- `/etc/init.d` – A directory also used by the SysVinit daemon, typically on Debian-based distributions, such as Ubuntu

Each script in `/etc/rc.d/init.d` takes care of all that is needed for starting or stopping a particular service on the server. The following is a partial example of the `httpd` script on a Linux system that uses the SysVinit daemon. It contains a case statement for handling the parameter (`$1`) that was passed to it, such as `start`, `stop`, `status`, and so on.

```
# cat /etc/rc.d/init.d/httpd
#!/bin/bash
#
# httpd          Startup script for the Apache HTTP Server
#
# chkconfig: - 85 15
# description: Apache is a World Wide Web server.
#              It is used to serve \
#              HTML files and CGI.
# processname: httpd
# config: /etc/httpd/conf/httpd.conf
# config: /etc/sysconfig/httpd
# pidfile: /var/run/httpd.pid

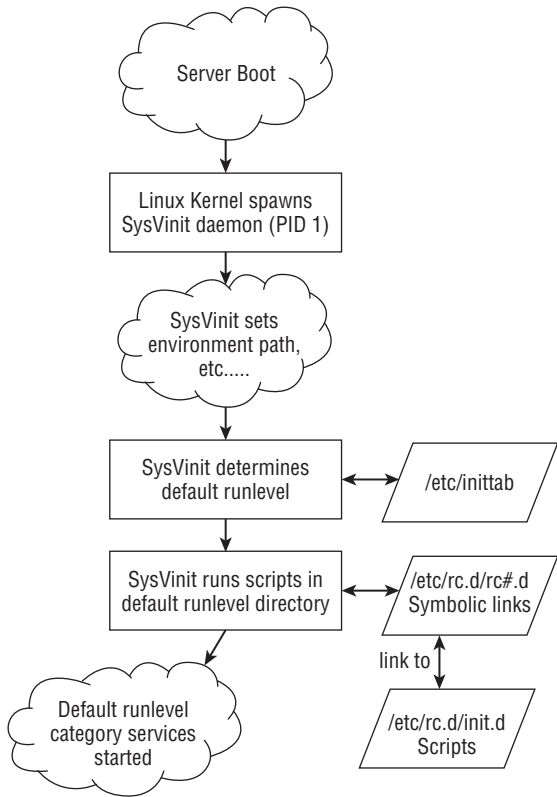
# Source function library.
. /etc/rc.d/init.d/functions
...
# See how we were called.
case "$1" in
    start)
        start
        ;;
    stop)
        stop
        ;;
    status)
        status $httpd
        RETVAL=$?
        ;;
    ...
esac

exit $RETVAL
```

Once the desired `runlevel` scripts in the `/etc/init.d` directory are executed, then the `SysVinit` daemon's process spawning is complete. This entire course is a well-organized, step-by-step approach for starting up a Linux server with the appropriate services running. Figure 15.1 reviews the `SysVinit` daemon's role in spawning the correct processes at system startup.

FIGURE 15.1

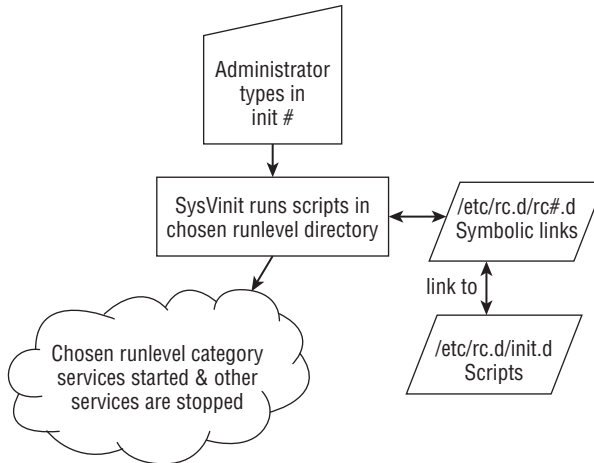
Classic `SysVinit` daemon course of action at server boot



The course is very similar when switching from one `runlevel` category to another. Figure 15.2 reviews the `SysVinit` daemon's roles in stopping and starting the correct processes during a `runlevel` switch.

FIGURE 15.2

Classic SysVinit daemon course of action at runlevel change



Now that you have a grasp of the classic `init` daemons, it's time to move on to the more modern `Upstart` `init` daemon.

Understanding the Upstart init daemon

As mentioned earlier, many Linux distributions have moved from the classic `init` daemons to the more modern `Upstart` `init` daemon. Included in that distribution list are the RHEL and Ubuntu distributions.

Learning Upstart init daemon basics

The primary difference between the classics and `Upstart` is the handling of stopping and starting services. The `SysVinit` daemon was created to operate in a static environment. The `Upstart` `init` daemon was created to operate in a flexible and ever-changing environment. With `SysVinit`, services are stopped and started based upon runlevels. The `Upstart` `init` daemon is not concerned with runlevels but with system events. Events are what determine when services are stopped and/or started. An *event* is a Linux server occurrence that triggers a needed system state change, which is communicated to the `Upstart` `init` daemon. The following are some examples of system events:

- The server boots up.
- The `init` command is used.
- A USB device is plugged into the server.

While the classic `init` daemons could handle the first two event examples, they could not deal well with the third.

Upstart handles services through defined jobs. An Upstart *job* can be either a task or a service. A *task* performs a limited duty, completes its work, and then returns to a waiting state. A *service*, on the other hand, is a long running program that never finishes its work or self-terminates, but instead stays in a running state. A daemon is an example of an Upstart service job.

The example that follows shows several Upstart jobs that include both task and service jobs. The task jobs are in a stop/waiting state, such as the task `rc`. The service jobs are in a start/running state, such as the cups daemon.

```
$ initctl list
avahi-daemon start/running, process 456
mountall-net stop/waiting
rc stop/waiting
rsyslog start/running, process 411
...
ssh start/running, process 405
udev-fallback-graphics stop/waiting
control-alt-delete stop/waiting
hwclock stop/waiting
mounted-proc stop/waiting
network-manager start/running, process 458
...
rc-sysinit stop/waiting
cups start/running, process 1066
...
tty6 start/running, process 833
ureadahead stop/waiting
```

These various jobs are defined via a jobs definition file. All the job definition files are located in the `/etc/init` directory as shown here:

```
$ ls /etc/init
acpid.conf          networking.conf
alsa-restore.conf   network-interface.conf
alsa-store.conf      network-interface-security.conf
anacron.conf         network-manager.conf
control-alt-delete.conf  procps.conf
cron.conf           rc.conf
cups.conf           rcS.conf
dbus.conf           rc-sysinit.conf
dmesg.conf          rsyslog.conf
failsafe.conf        setvtrgb.conf
friendly-recovery.conf  ssh.conf
hostname.conf        tty1.conf
hwclock.conf         tty2.conf
hwclock-save.conf    tty3.conf
```

```

irqbalance.conf      tty4.conf
lightdm.conf          tty5.conf
...

```

The Upstart `init` daemon depends upon events to trigger certain services to start, stop, restart, and so on. Events are either communicated to the Upstart `init` daemon or they are created by the Upstart daemon. This is called an *emitted event*. The actions taken when an event is emitted are dependent upon the settings in a job's configuration file. Consider the following Network Manager daemon's configuration file:

```

$ cat /etc/init/network-manager.conf
# network-manager - network connection manager
#
# The Network Manager daemon manages the system's network connections,
# automatically switching between the best available.

description      "network connection manager"

start on (local-filesystems      and started dbus)
stop on stopping dbus

expect fork
respawn

exec NetworkManager
$

```

From the example, you can see that there are two events that must take place in order to trigger the Upstart `init` daemon to start the Network Manager daemon:

- **The local-filesystems event** — The Upstart `init` daemon will emit this event when all the local filesystems in the `/etc/fstab` configuration file have been mounted.
- **The dbus daemon started event** — The Upstart `init` daemon will emit this started event when the `dbus` daemon has reached the `start/running` state.

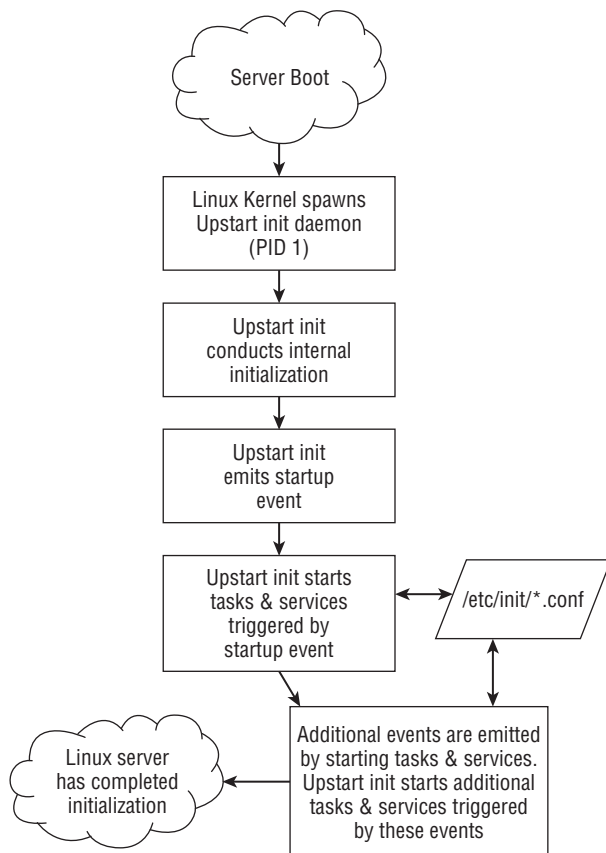
Thus, when these two events occur, the Upstart `init` daemon is informed and then will start the `NetworkManager` daemon.

Because the Upstart `init` daemon can handle these events and tracks the status (state) of processes, it is often referred to as a “state machine.” The Upstart `init` daemon is also referred to as an “event engine” because it emits events itself.

The Upstart `init` daemon is a flexible, organized, and clever approach for dealing with server events. Figure 15.3 shows the Upstart `init` daemon's course of action in spawning the correct processes at system startup time.

FIGURE 15.3

Upstart init daemon course of action at server boot



Learning Upstart's backward compatibility to SysVinit

Upstart provides backward compatibility to the SysVinit daemon. This has allowed the Linux distributions time to slowly migrate to Upstart.

The `/etc/inittab` file is still on some distributions. RHEL and the Fedora distributions still using Upstart use `/etc/inittab` to boot to the default runlevel listed. The Ubuntu distribution no longer has the file. The example of the `/etc/inittab` file that follows comes from a server running a version of Fedora, which uses the Upstart init daemon.

```
$ cat /etc/inittab
# inittab is only used by upstart for the default runlevel.
#
```

```
# ADDING OTHER CONFIGURATION HERE WILL HAVE NO EFFECT ON YOUR SYSTEM.
#
...
#
id:5:initdefault:
```

As you can see from the comment lines in the `/etc/inittab` file, the only thing this file is used for on Linux distributions that maintain it is to change the default runlevel at server boot time.

Tip

To change the default runlevel on an Ubuntu distribution that uses Upstart, edit `/etc/init/rc-sysinit.conf` and change the line `env DEFAULT_RUNLEVEL=#` where `#` is 2 to 5. However, remember that the runlevels 2-5 on Ubuntu are equivalent to SysVinit runlevel 5. Therefore, this activity is rather pointless.

System initialization compatibility to SysVinit is maintained on some distributions, such as Ubuntu, via the `/etc/init/rc-sysinit.conf` configuration file. This is one of the configuration files used at system boot, as shown previously in Figure 15.3. In the example that follows, you can see that Upstart checks for a `/etc/inittab` file and also runs any scripts that may still be in the `/etc/init.d/rcS` directory:

```
$ cat /etc/init/rc-sysinit.conf
# rc-sysinit - System V initialisation compatibility
#
# This task runs the old System V-style system initialisation scripts,
# and enters the default runlevel when finished.
...
start on (filesystem and static-network-up) or failsafe-boot
stop on runlevel

# Default runlevel, this may be overridden on the kernel command-line
# or by faking an old /etc/inittab entry
env DEFAULT_RUNLEVEL=2

emits runlevel
...
task

script
    # Check for default runlevel in /etc/inittab
    if [ -r /etc/inittab ]
    then
        eval "$(sed -nre 's/^[^#][^:]*:([0-6sS]):initdefault:
.*\/DEFAULT_RUNLEVEL="\1"/;p' /etc/inittab || true)"
    fi

    # Check kernel command-line for typical arguments
```

```
for ARG in $(cat /proc/cmdline)
do
    case "${ARG}" in
        -b|emergency)
            # Emergency shell
            [ -n "${FROM_SINGLE_USER_MODE}" ] || slogin
            ;;
        [0123456sS])
            # Override runlevel
            DEFAULT_RUNLEVEL="${ARG}"
            ;;
        -s|single)
            # Single user mode
            [ -n "${FROM_SINGLE_USER_MODE}" ] || DEFAULT_RUNLEVEL=S
            ;;
    esac
done

# Run the system initialisation scripts
[ -n "${FROM_SINGLE_USER_MODE}" ] || /etc/init.d/rcS

# Switch into the default runlevel
telinit "${DEFAULT_RUNLEVEL}"
end script
```

As you can see from the preceding example, the runlevel concept is maintained in the Upstart `init` daemon. In fact, there is even a runlevel signal that Upstart can emit.

```
# man -k "event signal"
control-alt-delete    (7) - event signalling console press of Control-Alt-Delete
keyboard-request      (7) - event signalling console press of Alt-UpArrow
power-status-changed  (7) - event signalling change of power status
runlevel              (7) - event signalling change of system runlevel
started               (7) - event signalling that a job is running
starting              (7) - event signalling that a job is starting
startup               (7) - event signalling system startup
stopped               (7) - event signalling that a job has stopped
stopping              (7) - event signalling that a job is stopping
```

Switching to a different runlevel is still allowed through the `init` or `telinit` commands. Any runlevel event is handled by the `rc` task.

```
$ initctl status rc
rc stop/waiting
```

The `rc` task job's configuration file is shown next. When a runlevel event is emitted, the `rc` configuration file will call the `/etc/rc.d/rc` script. When called, the `/etc/rc.d/rc` script will run the scripts located in the `/etc/rc.d/rc#.d`, where `#` is the chosen runlevel. This provides runlevel backward compatibility to SysVinit.


```

$ cat /etc/init/rc.conf
# rc - System V runlevel compatibility
#
# This task runs the old sysv-rc runlevel scripts. It
# is usually started by the telinit compatibility wrapper.

start on runlevel [0123456]

stop on runlevel [!$RUNLEVEL]

task

export RUNLEVEL
console output
exec /etc/rc.d/rc $RUNLEVEL

```

If you look back at the `/etc/inittab` in the classic SysVinit daemon section, you will notice that `/etc/inittab` also handled spawning the `getty` or `mingetty` processes. The Upstart `init` daemon handles this via the `start-ttys` task.

```

# initctl status start-ttys
start-ttys stop/waiting

```

The `start-ttys` task job's configuration file is shown next. When a runlevel event is emitted, the `start-ttys` configuration file will spawn the `getty` or `mingetty` process.

```

$ cat /etc/init/start-ttys.conf
#
# This service starts the configured number of gettys.

start on stopped rc RUNLEVEL=[2345]

env ACTIVE_CONSOLES=/dev/tty[1-6]
env X_TTY=/dev/tty1
task
script
. /etc/sysconfig/init
for tty in $(echo $ACTIVE_CONSOLES) ; do
    [ "$RUNLEVEL" = "5" -a "$tty" = "$X_TTY" ] && continue
    initctl start tty TTY=$tty
done
end script

```

Although the Upstart `init` daemon provides backward compatibility to the classic SysVinit daemon, is a state-machine, and can handle ever-changing events on a server, it is not the only modern init daemon available for the Linux server. Another modern init daemon is `systemd`.

Understanding systemd init

The `systemd` init daemon is also called `systemd` or the `system` daemon. `systemd` can replace the `SysVinit` and the `Upstart` init daemons. This modern init daemon currently runs on Fedora 15 and above and is backward compatible with both `SysVinit` and `Upstart`. System initialization time is reduced by `systemd` because it starts fewer services and starts them in a parallel manner. In addition, `systemd` can handle a fluid environment because it supervises all the processes on the entire Linux server.

Learning systemd basics

With the `SysVinit` daemon, services are stopped and started based upon runlevels. The `systemd` is also concerned with runlevels, but they are called `target` units. Units are the focus of `systemd`. A *unit* is a group consisting of a name, type, and configuration file and is focused on a particular service or action. The eight `systemd` units types are:

- `automount`
- `device`
- `mount`
- `path`
- `service`
- `snapshot`
- `socket`
- `target`

The two primary `systemd` units you need to be concerned with for dealing with services are `service` units and `target` units. A *service unit* is for managing daemons on your Linux server. A *target unit* is simply a group of other units.

The example that follows shows several `systemd` `service` units and `target` units. The `service` units have familiar daemon names, such as `cups` and `sshd`. Note that each `service` unit name ends with `.service`. The `target` units shown have names like `sysinit`. (`sysinit` is used for starting up services at system initialization.) The `target` unit names end with `.target`.

```
# systemctl list-units | grep .service
...
cups.service          loaded active running CUPS Printing Service
dbus.service          loaded active running D-Bus System Message Bus
...
NetworkManager.service loaded active running Network Manager
prefdm.service         loaded active running Display Manager
remount-rootfs.service loaded active exited Remount Root FS
rsyslog.service        loaded active running System Logging Service
...
```

```

sshd.service          loaded active running OpenSSH server daemon
system-s...yboard.service loaded active running System Setup Keyboard
systemd-logind.service loaded active running Login Service
...
#
# systemctl list-units | grep .target
basic.target          loaded active active Basic System
cryptsetup.target     loaded active active Encrypted Volumes
getty.target          loaded active active Login Prompts
graphical.target      loaded active active Graphical Interface
local-fs-pre.target   loaded active active Local File Systems (Pre)
local-fs.target        loaded active active Local File Systems
multi-user.target      loaded active active Multi-User
network.target        loaded active active Network
remote-fs.target      loaded active active Remote File Systems
sockets.target        loaded active active Sockets
sound.target          loaded active active Sound Card
swap.target           loaded active active Swap
sysinit.target        loaded active active System Initialization
syslog.target         loaded active active Syslog

```

The Linux system unit configuration files are located in the `/lib/systemd/system` and `/etc/systemd/system` directories. You could use the `ls` command to look through those directories, but the preferred method is to use an option on the `systemctl` command as follows:

```

# systemctl list-unit-files --type=service
UNIT FILE                                STATE
...
cups.service                            enabled
...
dbus.service                            static
...
NetworkManager.service                 enabled
...
poweroff.service                        static
...
sshd.service                            enabled
sssd.service                            disabled
...
134 unit files listed.

```

The unit configuration files shown in the preceding code are all associated with a service unit. Configuration files for Target units can be displayed via the following method.

```

# systemctl list-unit-files --type=target
UNIT FILE                                STATE
anaconda.target                         static
basic.target                            static

```

```
bluetooth.target          static
cryptsetup.target         static
ctrl-alt-del.target       disabled
default.target            enabled
...
shutdown.target          static
sigpwr.target             static
smartcard.target         static
sockets.target           static
sound.target             static
swap.target              static
sysinit.target           static
syslog.target            static
time-sync.target         static
umount.target            static
43 unit files listed.
```

Notice that both of the configuration units' file examples shown display units with a status of either static, enabled, or disabled. The enabled status means that the unit is currently enabled. The disabled status means that the unit is currently disabled. The next status, static, is slightly confusing. It stands for "statically enabled," and it means that the unit is enabled by default and cannot be disabled, even by root.

The service unit configuration files contain a lot of information, such as what other services must be started, when this service can be started, which environmental file to use, and so on. The following example shows the `sshd`'s unit configuration file:

```
# cat /lib/systemd/system/sshd.service
[Unit]
Description=OpenSSH server daemon
After=syslog.target network.target auditd.service

[Service]
EnvironmentFile=/etc/sysconfig/sshd
ExecStart=/usr/sbin/sshd -D $OPTIONS
ExecReload=/bin/kill -HUP $MAINPID

[Install]
WantedBy=multi-user.target
#
```

This basic service unit configuration file has the following options:

- **Description** — This is a free-form description (comment line) of the service.
- **After** — This setting configures ordering. In other words, it lists which units should be activated before this service is started.
- **Environment File** — The service's configuration file.
- **ExecStart** — The command used to start this service.

- **ExecReload** — The command used to reload this service.
- **WantedBy** — This identifies what `target` unit this service belongs to.

Notice that the `target` unit, `multi-user.target`, is used in the `sshd` service unit configuration file. The `sshd` service unit is wanted by the `multi-user.target`. In other words, when the `multi-user.target` unit is activated, the `sshd` service unit will be started.

You can view the various units a `target` unit will activate by using the following command:

```
# systemctl show --property "Wants" multi-user.target
Wants=multipathd.service avahi-daemon.service sshd-keygen.se
(END) q
```

Unfortunately, the `systemctl` command does not format the output for this well. It literally runs off the right edge of the screen so that you cannot see the full results. And you have to enter a `q` to return to the command prompt. To fix this problem, pipe the output through some formatting commands to produce a nice alphabetically sorted display, as shown in the example that follows.

```
# systemctl show --property "Wants" multi-user.target |
> fmt -10 | sed 's/Wants=//g' | sort
abrt-ccpp.service
abrt-d.service
abrt-oops.service
abrt-vmcore.service
atd.service
auditd.service
avahi-daemon.service
crond.service
cups.path
dbus.service
fcoe.service
getty.target
irqbalance.service
iscsid.service
iscsi.service
livesys-late.service
livesys.service
lldpad.service
mcelog.service
mdmonitor.service
multipathd.service
netfs.service
NetworkManager.service
plymouth-quit.service
plymouth-quit-wait.service
remote-fs.target
```

```
rsyslog.service
sendmail.service
sm-client.service
sshd-keygen.service
sshd.service
systemd-ask-password-wall.path
systemd-logind.service
systemd-update-utmp-runlevel.service
systemd-user-sessions.service
#
```

This display shows all the services and various other units that will be activated (started), including the `sshd`, when the `multi-user.target` unit is activated. Remember that a `target` unit is simply a grouping of other units, as shown in the preceding example. Also notice that the units in this group are not all `service` units. There are `path` units and another `target` unit as well.

A `target` unit has both `Wants` and requirements, called `Requires`. A *Wants* means that all the units listed are triggered to activate (start). If they fail or cannot be started, no problem — the `target` unit continues on its merry way. The preceding example is a display of `Wants` only.

A `Requires` is much more stringent and potentially catastrophic than a `Wants`. A *Requires* means that all the units listed are triggered to activate (start). If they fail or cannot be started, the entire unit (group of units) is deactivated.

You can view the various units a `target` unit `Requires` (must activate or the unit will fail), using the command in the example that follows. Notice the `Requires` output is much shorter than the `Wants` for the `multi-user.target`. Thus, no special formatting of the output is needed.

```
# systemctl show --property "Requires" multi-user.target
Requires=basic.target
```

The `target` units also have configuration files, as do the `service` units. The following example shows the contents of the `multi-user.target` configuration file.

```
# cat /lib/systemd/system/multi-user.target
# This file is part of systemd.
#
...

[Unit]
Description=Multi-User
Requires=basic.target
Conflicts=rescue.service rescue.target
```

```

After=basic.target rescue.service rescue.target
AllowIsolate=yes

[Install]
Alias=default.target

```

This basic `target` unit configuration file has the following options:

- **Description** — This is just a free-form description of the target.
- **Requires** — If this `multi-user.target` gets activated, then the listed `target` unit will also be activated. If the listed `target` unit is deactivated or fails, then `multi-user.target` will be deactivated. If there are no `After` and `Before` options, then both `multi-user.target` and listed `target` unit will activate simultaneously.
- **Conflicts** — This setting avoids conflicts in services. Starting `multi-user.target` will stop the listed targets and services, and vice-versa.
- **After** — This setting configures ordering. In other words, it determines which units should be activated before starting this service.
- **AllowIsolate** — This option is a Boolean setting of yes or no. If set to yes, then this `target` unit, `multi-user.target`, is activated along with its dependencies and all others are deactivated.
- **ExecStart** — This command starts the service.
- **ExecReload** — This command reloads the service.
- **Alias** — With this command, `systemd` will create a symbolic link from the `target` unit names listed to this unit, `multi-user.target`.

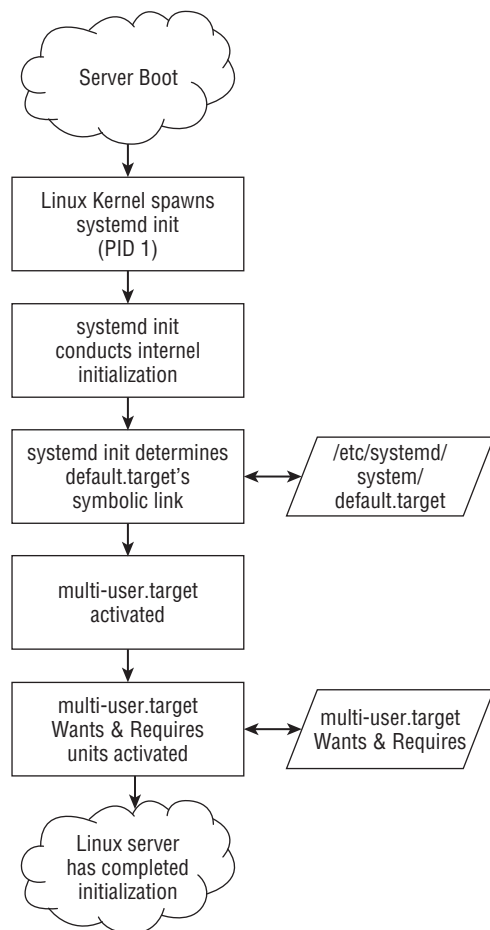
To get more information on these configuration files and their options, enter **man systemd.service**, **man systemd.target**, and **man systemd.unit** at the command line.

For the Linux server using `systemd` init, the boot process will be easier to follow, now that you understand `systemd` `target` units. At boot, `systemd` activates the `default.target` unit. This unit is aliased to either `multi-user.target` or `graphical.target`. Thus, depending upon the alias set, the services targeted by the `target` unit are started. Figure 15.4 shows the `systemd` init's course of action in spawning the correct processes at system startup. If you are following along with the examples used in this section, the `default.target` is aliased to `multi-user.target`, as shown in Figure 15.4.

If you need more help on understanding `systemd` init, there is excellent documentation by its author at <http://0pointer.de/blog/projects/systemd-docs.html>. Also, you can enter **man -k systemd** at the command line to get a listing of the various `systemd` utilities' documentation in the man pages.

FIGURE 15.4

systemd init course of action at server boot



Learning systemd's backward compatibility to SysVinit

The `systemd init` has maintained backward compatibility to the `SysVinit` daemon. This allows Linux distributions time to slowly migrate to `systemd`.

While runlevels are not truly part of `systemd`, the `systemd` infrastructure has been created to provide compatibility with the concept of runlevels. There are seven target units configuration files specifically created for backward compatibility to `SysVinit`:

- `runlevel0.target`
- `runlevel1.target`

- `runlevel2.target`
- `runlevel3.target`
- `runlevel4.target`
- `runlevel5.target`
- `runlevel6.target`

As you probably have already figured out, there is a `target` unit configuration file for each of the seven classic SysVinit runlevels. These `target` unit configuration files are symbolically linked to `target` unit configuration files that most closely match the idea of the original runlevel. In the example that follows, the symbolic links are shown for `runlevel` target units. Notice that the `runlevel` target units for runlevel 2, 3, and 4 are all symbolically linked to `multi-user.target`. The `multi-user.target` unit is similar to the legacy Extended Multi-user Mode.

```
# ls -l /lib/systemd/system/runlevel*.target
lrwxrwxrwx. 1 root root 15 Mar 27 15:39
  /lib/systemd/system/runlevel0.target -> poweroff.target
lrwxrwxrwx. 1 root root 13 Mar 27 15:39
  /lib/systemd/system/runlevel1.target -> rescue.target
lrwxrwxrwx. 1 root root 17 Mar 27 15:39
  /lib/systemd/system/runlevel2.target -> multi-user.target
lrwxrwxrwx. 1 root root 17 Mar 27 15:39
  /lib/systemd/system/runlevel3.target -> multi-user.target
lrwxrwxrwx. 1 root root 17 Mar 27 15:39
  /lib/systemd/system/runlevel4.target -> multi-user.target
lrwxrwxrwx. 1 root root 16 Mar 27 15:39
  /lib/systemd/system/runlevel5.target -> graphical.target
lrwxrwxrwx. 1 root root 13 Mar 27 15:39
  /lib/systemd/system/runlevel6.target -> reboot.target
#
```

The `/etc/inittab` file still exists, but it only contains comments stating this configuration file is not used and gives some basic `systemd` information. The `/etc/inittab` file no longer has any true functional use. The following is an example of a `/etc/inittab` file on a Linux server that uses `systemd`.

```
# cat /etc/inittab
# inittab is no longer used when using systemd.
#
# ADDING CONFIGURATION HERE WILL HAVE NO EFFECT ON YOUR SYSTEM.
#
# Ctrl-Alt-Delete is handled by
/etc/systemd/system/ctrl-alt-del.target
#
# systemd uses 'targets' instead of runlevels.
  By default, there are two main targets:
#
# multi-user.target: analogous to runlevel 3
```

```
# graphical.target: analogous to runlevel 5
#
# To set a default target, run:
#
# ln -s /lib/systemd/system/<target name>.target
# /etc/systemd/system/default.target
```

The `/etc/inittab` explains that if you want something similar to a classic 3 or 5 runlevel as your default runlevel, you will need to create a symbolic link from the `default.target` unit to the runlevel target unit of your choice. To check what `default.target` is currently symbolically linked to (or in legacy terms, to check the default runlevel), use the command shown here. You can see that on this Linux server, the default is to start up at legacy runlevel 3.

```
# ls -l /etc/systemd/system/default.target
lrwxrwxrwx. 1 root root 36 Mar 13 17:27
/etc/systemd/system/default.target ->
/lib/systemd/system/runlevel3.target
```

The capability to switch runlevels using the `init` or `telinit` command is still available. When issued, either of the commands will be translated into a `systemd` target unit activation request. Therefore, typing `init 3` at the command line really issues the command `systemctl isolate multi-user.target`. Also, you can still use the `runlevel` command to determine the current legacy runlevel, but it is strongly discouraged.

The classic `SysVinit` `/etc/inittab` handled spawning the `getty` or `mingetty` processes. The `systemd` `init` handles this via the `getty.target` unit. The `getty.target` is activated by the `multi-user.target` unit. You can see how these two target units are linked by the following command:

```
# systemctl show --property "WantedBy" getty.target
WantedBy=multi-user.target
```

As of Fedora 16, the conversion of `SysVinit` core services to `systemd` is complete. However, to find out the current status of a particular service, visit the `systemd` compatibility web page at <http://fedoraproject.org/wiki/User:Johannbg/QA/Systemd/compatibility>.

Now that you have a basic understanding of classic and modern `init` daemons, it's time to look at doing some practical server administrator actions that involve the `init` daemon.

Auditing Services

As a Linux administrator, you need to audit the services being offered on your server for documentation, security, and troubleshooting purposes. Typically, you will audit services for documentation purposes if you have “inherited” a server, although some sites

do this on a regular basis. For security reasons, you will need to disable and remove any unused system services discovered through the audit process. Most importantly for troubleshooting purposes, you need to be able to quickly know what should and should not be running on your Linux server.

Of course, knowing which `init` daemon is being used by your Linux server is the first piece of information to obtain. How to determine this was covered in the “Understanding the Linux `init` Daemon” section of this chapter. The rest of this section is organized into subsections on the various `init` daemons.

Auditing the classic SysVinit daemon

To see all the services that are being offered by a Linux server using the classic SysVinit daemon, use the `chkconfig` command. The example that follows shows the services available on a classic SysVinit Linux server. Note that each runlevel (0–6) is shown for each service with a status of `on` or `off`. The status denotes whether a particular service is started (`on`) or not (`off`) for that runlevel.

```
# chkconfig --list
ConsoleKit      0:off  1:off  2:off  3:on   4:on   5:on   6:off
NetworkManager 0:off  1:off  2:off  3:off  4:off  5:off  6:off
...
crond           0:off  1:off  2:on   3:on   4:on   5:on   6:off
cups          0:off  1:off  2:on   3:on   4:on   5:on   6:off
...
sshd            0:off  1:off  2:on   3:on   4:on   5:on   6:off
syslog          0:off  1:off  2:on   3:on   4:on   5:on   6:off
tux             0:off  1:off  2:off  3:off  4:off  5:off  6:off
udev-post       0:off  1:off  2:off  3:on   4:on   5:on   6:off
vncserver     0:off  1:off  2:off  3:off  4:off  5:off  6:off
winbind         0:off  1:off  2:off  3:off  4:off  5:off  6:off
wpa_supplicant  0:off  1:off  2:off  3:off  4:off  5:off  6:off
xfs             0:off  1:off  2:on   3:on   4:on   5:on   6:off
ypbind          0:off  1:off  2:off  3:off  4:off  5:off  6:off
yum-updatesd    0:off  1:off  2:off  3:on   4:on   5:on   6:off
```

Some services in the example are never started, such as `vncserver`. Other services, such as the `cups` daemon, are started on runlevels 2 through 5.

Using the `chkconfig` command, you cannot tell if a service is currently running. To do that, you will need to use the `service` command. To help isolate only those services that are currently running, the `service` command is piped into the `grep` command and then sorted, as follows.

```
# service --status-all | grep running... | sort
anacron (pid 2162) is running...
atd (pid 2172) is running...
```

```
auditd (pid 1653) is running...
automount (pid 1952) is running...
console-kit-daemon (pid 2046) is running...
crond (pid 2118) is running...
cupsd (pid 1988) is running...
...
sshd (pid 2002) is running...
syslogd (pid 1681) is running...
xfs (pid 2151) is running...
yum-updatesd (pid 2205) is running...
```

You can also use both the `chkconfig` and the `service` commands to view an individual service's settings. Using both commands in the example that follows, you can view the cups daemon's settings.

```
# chkconfig --list cups
cups          0:off  1:off  2:on   3:on   4:on   5:on   6:off
#
# service cups status
cupsd (pid 1988) is running...
```

You can see that cupsd is set to start on every runlevel but 0, 1, and 6, and from the `service` command, you can see that it is currently running. Also, the process ID (PID) number is given for the daemon.

Auditing the Upstart init daemon

To see all the services running on a Linux server using the Upstart `init` daemon, use the following command:

```
# initctl list | grep start/running
tty (/dev/tty3) start/running, process 1163
...
system-setup-keyboard start/running, process 656
prefdm start/running, process 1154
```

Keep in mind that many services may have not yet been ported to the Upstart `init` daemon. Therefore, you will also need to use the classic SysVinit command, `service`, to check for any leftover SysVinit services. Note that on some distributions, you may see a few services in *both* the `initctl` and the `service` command output.

```
# service --status-all | grep running
abrttd (pid 1118) is running...
acpid (pid 996) is running...
atd (pid 1146) is running...
...
rsyslogd (pid 752) is running...
sendmail (pid 1099) is running...
...
```

CAUTION

Just because a service is not in a running state does not mean it is unavailable. The service could be in a stopped/wait state, awaiting an event on the system. To see all the services, no matter what their state, remove the `grep` portion of the preceding `initctl list` and `service --status-all` commands.

To show the status of a single service, use `initctl` if the service has been ported to Upstart and the `service` command if it has not been ported yet. The following example shows two service statuses — one that has been ported to Upstart and one that has not.

```
# initctl status vpnc-cleanup
vpnc-cleanup stop/waiting
#
# service ssh status
sshd (pid 970) is running...
```

The `ssh` daemon has not yet been ported, on this Linux server, to Upstart. Therefore, `ssh` needs the `service` command with the `status` option to be used for auditing. The `vpnc-cleanup` service is an Upstart service. Thus, it needed the `initctl status` command to be used. In some distributions, such as Ubuntu, you can also use the `initctl status` command for services that have not yet been migrated to Upstart.

Auditing the systemd init

To see all the services that are being offered by a Linux server using `systemd`, use the following command:

```
# systemctl list-unit-files --type=service | grep -v disabled
UNIT FILE                                STATE
abrt-ccpp.service                       enabled
abrt-oops.service                       enabled
abrt-vmcore.service                     enabled
abrtd.service                           enabled
alsa-restore.service                    static
alsa-store.service                      static
anaconda-shell@.service                  static
arp-ethers.service                      enabled
atd.service                             enabled
auditd.service                          enabled
avahi-daemon.service                    enabled
bluetooth.service                       enabled
console-kit-log-system-restart.service  static
console-kit-log-system-start.service    static
console-kit-log-system-stop.service     static
crond.service                           enabled
cups.service                            enabled
...
```

```
sshd-keygen.service          enabled
sshd.service                 enabled
system-setup-keyboard.service enabled
...
134 unit files listed.
```

Remember that the three status possibilities for a `systemd` service are `enabled`, `disabled`, or `static`. There's no need to include `disabled` in this audit, which is effectively accomplished by using the `-v` option on the `grep` command, as shown in the preceding example. The state of `static` is essentially `enabled`, and thus should be included.

NOTE

Most of the system services have been ported to `systemd`. However, if you want to double check, you can see the status of a particular service's migration at <http://fedoraproject.org/wiki/User:Johannbg/QA/Systemd/compatibility>.

To see if a particular service is running, use the following command:

```
# systemctl status cups.service
cups.service - CUPS Printing Service
Loaded: loaded (/lib/systemd/system/cups.service; enabled)
Active: active (running) since Mon, 30 Apr 2015 12:36:31 -0400; 13h ago
Main PID: 1315 (cupsd)
CGroup: name=systemd:/system/cups.service
        └ 1315 /usr/sbin/cupsd -f
```

The `systemctl` command can be used to show the status of a single service. In the preceding example, the printing service was chosen. Notice that the name of the service is `cups.service`. A great deal of helpful information about the service is given here, such as the fact that it is `enabled` and `active`, its start time, and its process ID (PID) as well.

Now that you can audit services and determine some information about them, you need to know how to accomplish starting, stopping, and reloading the services on your Linux server.

Stopping and Starting Services

The tasks of starting, stopping, and restarting services typically refer to immediate needs — in other words, managing services without a server reboot. For example, if you want to temporarily stop a service, then you are in the right section. However, if you want to stop a service and not allow it to be restarted at server reboot, then you need to actually disable the service, which is covered in the section “Configuring Persistent Services,” later in this chapter.

Stopping and starting the classic SysVinit daemon

The primary command for stopping and starting SysVinit services is the `service` command. With the `service` command, the name of the service you are wishing to control comes second in the command line. The last option is what you want to do to the service, `stop`, `start`, `restart`, and so on. The following example shows how to stop the `cups` daemon. Notice that an OK is given, which lets you know that `cupsd` has been successfully stopped.

```
# service cups status
cupsd (pid 5857) is running...
#
# service cups stop
Stopping cups:          [ OK ]
#
# service cups status
cupsd is stopped
```

To start a service, you simply add a `start` option instead of a `stop` option on the end of the `service` command as follows.

```
# service cups start
Starting cups:          [ OK ]
#
# service cups status
cupsd (pid 6860) is running...
```

To restart a SysVinit service, the `restart` option is used. This option will stop the service and then immediately start it again.

```
# service cups restart
Stopping cups:          [ OK ]
Starting cups:          [ OK ]
#
# service cups status
cupsd (pid 7955) is running...
```

When a service is already stopped, a `restart` will generate a `FAILED` status on the attempt to stop it. However, as shown in the example that follows, the service will be successfully started when a `restart` is attempted.

```
# service cups stop
Stopping cups:          [ OK ]
#
# service cups restart
Stopping cups:          [FAILED]
Starting cups:          [ OK ]
#
# service cups status
cupsd (pid 8236) is running...
```

Reloading a service is different from restarting a service. When you reload a service, the service itself is not stopped. Only the service's configuration files are loaded again. The following example shows how to reload the cups daemon.

```
# service cups status
cupsd (pid 8236) is running...
#
# service cups reload
Reloading cups:                [ OK ]
#
# service cups status
cupsd (pid 8236) is running...
```

If a SysVinit service is stopped when you attempt to reload it, you will get a FAILED status. This is shown in the following example:

```
# service cups status
cupsd is stopped
#
# service cups reload
Reloading cups:                [FAILED]
```

Stopping and starting the Upstart init daemon

The primary command for stopping and starting Upstart init services is the `initctl` command. The options are very similar to SysVinit's `service` command:

- **Stopping a service with Upstart init** — In the following example, the status of the cups daemon is checked and then stopped using the `initctl stop cups.service` command.

```
# initctl status cups
cups start/running, process 2390
#
# initctl stop cups
cups stop/waiting
#
# initctl status cups
cups stop/waiting
```

- **Starting a service with Upstart init** — In the following example, the cups daemon is stopped using the `initctl stop cups.service` command.

```
# initctl start cups
cups start/running, process 2408
#
# initctl status cups
cups start/running, process 2408
```


- **Restarting a service with Upstart init** — Restarting a service with Upstart init will stop and then start the service. However, the configuration file will not be reloaded.

```
# initctl restart cups
cups start/running, process 2490
#
# initctl status cups
cups start/running, process 2490
#
```

- **Reloading a service with Upstart init** — Reloading will *not* stop and start the service. It only loads the configuration file again. This is the option to use when you have made changes to the configuration file.

The example below illustrates how to reload the cups daemon with initctl. Notice that the process ID (PID) is still 2490, which is the same as it was in the example for restarting the cups daemon because the process was not stopped and started in the reload process.

```
# initctl reload cups
#
# initctl status cups
cups start/running, process 2490
```

NOTE

You need root privileges to stop and start services. However, you do not need root privileges to check a service's status.

Stopping and starting the systemd daemon

For the systemd daemon, the systemctl command will work for stopping, starting, reloading, and restarting. The options to the systemctl command should look familiar.

Stopping a service with systemd

In the example that follows, the status of the cups daemon is checked and then stopped using the systemctl stop cups.service command:

```
# systemctl status cups.service
cups.service - CUPS Printing Service
   Loaded: loaded (/lib/systemd/system/cups.service; enabled)
   Active: active (running) since Mon, 30 Apr 2015 12:36:3...
   Main PID: 1315 (cupsd)
     CGroup: name=systemd:/system/cups.service
             └─ 1315 /usr/sbin/cupsd -f
#
# systemctl stop cups.service
```

```
#
# systemctl status cups.service
cups.service - CUPS Printing Service
   Loaded: loaded (/lib/systemd/system/cups.service; enabled)
   Active: inactive (dead) since Tue, 01 May 2015 04:43:4...
   Process: 1315 ExecStart=/usr/sbin/cupsd -f
   (code=exited, status=0/SUCCESS)
   CGroup: name=systemd:/system/cups.service
```

Notice that when the status is taken, after stopping the cups daemon, the service is **inactive (dead)** but still considered **enabled**. This means that the cups daemon will still be started upon server boot.

Starting a service with systemd

Starting the cups daemon is just as easy as stopping it. The example that follows demonstrates this ease.

```
# systemctl start cups.service
#
# systemctl status cups.service
cups.service - CUPS Printing Service
   Loaded: loaded (/lib/systemd/system/cups.service; enabled)
   Active: active (running) since Tue, 01 May 2015 04:43:5...
   Main PID: 17003 (cupsd)
   CGroup: name=systemd:/system/cups.service
           └─ 17003 /usr/sbin/cupsd -f
```

After the cups daemon is started, using `systemctl` with the status option shows the service is **active (running)**. Also, its process ID (PID) number, 17003, is shown.

Restarting a service with systemd

Restarting a service means that a service is stopped and then started again. If the service was not currently running, restarting it will simply start the service.

```
# systemctl restart cups.service
#
# systemctl status cups.service
cups.service - CUPS Printing Service
   Loaded: loaded (/lib/systemd/system/cups.service; enabled)
   Active: active (running) since Tue, 01 May 2015 04:45:2...
   Main PID: 17015 (cupsd)
   CGroup: name=systemd:/system/cups.service
           └─ 17015 /usr/sbin/cupsd -f
```

You can also perform a conditional restart of a service using `systemctl`. A conditional restart only restarts a service if it is currently running. Any service in an inactive state will not be started.

```
# systemctl status cups.service
cups.service - CUPS Printing Service
   Loaded: loaded (/lib/systemd/system/cups.service; enabled)
   Active: inactive (dead) since Tue, 01 May 2015 06:03:32...
   Process: 17108 ExecStart=/usr/sbin/cupsd -f
            (code=exited, status=0/SUCCESS)
   CGroup: name=systemd:/system/cups.service
#
# systemctl condrestart cups.service
#
# systemctl status cups.service
cups.service - CUPS Printing Service
   Loaded: loaded (/lib/systemd/system/cups.service; enabled)
   Active: inactive (dead) since Tue, 01 May 2015 06:03:32...
   Process: 17108 ExecStart=/usr/sbin/cupsd -f
            (code=exited, status=0/SUCCESS)
   CGroup: name=systemd:/system/cups.service
```

Notice in the example that the cups daemon was in an inactive state. When the conditional restart was issued, no error messages were generated! The cups daemon was not started because conditional restarts will affect active services. Thus, it is always a good practice to check the status of a service, after stopping, starting, conditionally restarting, and so on.

Reloading a service with systemd

Reloading a service is different from restarting a service. When you reload a service, the service itself is not stopped. Only the service's configuration files are loaded again.

```
# systemctl reload cups.service
Failed to issue method call: Job type reload is
not applicable for unit cups.service.
#
# systemctl reload sshd.service
#
# systemctl status sshd.service
sshd.service - OpenSSH server daemon
   Loaded: loaded (/lib/systemd/system/sshd.service; enabled)
   Active: active (running) since Mon, 30 Apr 2015 12:35:2...
   Process: 17009 ExecReload=/bin/kill -HUP $MAINPID
            (code=exited, status=0/SUCCESS)
   Main PID: 786 (sshd)
   CGroup: name=systemd:/system/sshd.service
            └─ 786 /usr/sbin/sshd -D
```

Doing a reload of a service, instead of a restart, will prevent any pending service operations from being aborted. A reload is a better method for a busy Linux server.

Now that you know how to stop and start services for troubleshooting and emergency purposes, you can learn how to enable and disable services.

Configuring Persistent Services

You use `stop` and `start` for immediate needs, not for services that need to be persistent. A *persistent* service is one that is started at server boot time. Services that need to be set as persistent are typically new services that the Linux server will be offering.

Configuring the classic SysVinit daemon persistent services

One of the nice features of the classic SysVinit daemon is that making a particular service persistent or removing its persistence is very easy to do. Consider the following example:

```
# chkconfig --list cups
cups          0:off 1:off 2:off 3:off 4:off 5:off 6:off
```

On this Linux server, the `cups` daemon is not started at any runlevel, as shown with the `chkconfig` command. You can also check and see if there are any start (S) symbol links set up in each of the seven runlevel directories, `/etc/rc.d/rc?.d`. Remember that SysVinit keeps symbolic links here for starting and stopping various services at certain runlevels. Each directory represents a particular runlevel; for example, `rc5.d` is for runlevel 5. Notice that only files starting with a `K` are listed, so there are links for killing off the `cups` daemon. None are listed with `S`, which is consistent with `chkconfig` that the `cups` daemon does not start at any runlevel on this server.

```
# ls /etc/rc.d/rc?.d/*cups
/etc/rc.d/rc0.d/K10cups  /etc/rc.d/rc3.d/K10cups
/etc/rc.d/rc1.d/K10cups  /etc/rc.d/rc4.d/K10cups
/etc/rc.d/rc2.d/K10cups  /etc/rc.d/rc5.d/K10cups
/etc/rc.d/rc6.d/K10cups
```

To make a service persistent at a particular runlevel, the `chkconfig` command is used again. Instead of the `--list` option, the `--level` option is used, as shown in the following code:

```
# chkconfig --level 3 cups on
#
# chkconfig --list cups
cups          0:off 1:off 2:off 3:on  4:off 5:off 6:off
#
# ls /etc/rc.d/rc3.d/S*cups
/etc/rc.d/rc3.d/S56cups
```

The service's persistence at runlevel 3 is verified by using both the `chkconfig --list` command and looking at the `rc3.d` directory for any files starting with the letter `S`.

To make a service persistent on more than one runlevel, you can do the following:

```
# chkconfig --level 2345 cups on
#
```

```
# chkconfig --list cups
cups          0:off 1:off 2:on  3:on  4:on  5:on  6:off
#
# ls /etc/rc.d/rc?.d/S*cups
/etc/rc.d/rc2.d/S56cups  /etc/rc.d/rc4.d/S56cups
/etc/rc.d/rc3.d/S56cups  /etc/rc.d/rc5.d/S56cups
```

Disabling a service is just as easy as enabling one with SysVinit. You just need to change the on in the `chkconfig` command to `off`. The following example demonstrates using the `chkconfig` command to disable the `cups` service at runlevel 5.

```
# chkconfig --level 5 cups off
#
# chkconfig --list cups
cups          0:off 1:off 2:on  3:on  4:on  5:off  6:off
#
# ls /etc/rc.d/rc5.d/S*cups
ls: cannot access /etc/rc.d/rc5.d/S*cups: No such file or directory
```

As expected, there is now no symbolic link, starting with the letter S, for the `cups` daemon in the `/etc/rc.d/rc5.d` directory.

Configuring Upstart init daemon persistent services

The Upstart `init` daemon emits the startup signal that triggers the service jobs to start. At server boot time, various jobs may themselves emit signals. These emitted signals then cause other jobs to start. Thus, the key to making a service persistent is to ensure the service's definition file is triggered by one of the signals emitted as the server boots.

Remember that the Upstart `init` daemon's job definition files are located in `/etc/init`. Consider the following job definition file for the `ssh` daemon:

```
# cat /etc/init/ssh.conf
# ssh - OpenBSD Secure Shell server
#
# The OpenSSH server provides secure shell access to the system.

description    "OpenSSH server"

start on filesystem or runlevel [2345]
stop on runlevel [!2345]

respawn
```

To determine what emitted event(s) trigger a service, look for `start on` in the configuration file. The `ssh` daemon is triggered by several possible emitted events, `filesystem`, `runlevel 2`, `runlevel 3`, `runlevel 4`, or `runlevel 5`. Basically, the `ssh` daemon starts upon server boot and is set as persistent. The syntax for the `runlevel` events, `runlevel [2345]`, is used in many of the job files and denotes that the name "runlevel" can end in either 2, 3, 4, or 5.

To make a job persistent (start at boot), you will need to modify the `start` on line in its configuration file so that it starts on certain events emitted at server boot. To disable a job at boot, just comment out the `start` on line with a pound sign (#). See the section “Adding New or Customized Services” for Upstart for a more thorough explanation of these configuration files.

Configuring systemd init persistent services

For the `systemd` daemon, again the `systemctl` command is used. With it, you can disable and enable services on the Linux server.

Enabling a service with systemd

Using the `enable` option on the `systemctl` command will set a service to always start at boot (be persistent). The following shows exactly how to accomplish this:

```
# systemctl status cups.service
cups.service - CUPS Printing Service
   Loaded: loaded (/lib/systemd/system/cups.service; disabled)
   Active: inactive (dead) since Tue, 01 May 2015 06:42:38 ...
   Main PID: 17172 (code=exited, status=0/SUCCESS)
   CGroup: name=systemd:/system/cups.service

#
# systemctl enable cups.service
ln -s '/lib/systemd/system/cups.service'
   '/etc/systemd/system/printer.target.wants/cups.service
ln -s '/lib/systemd/system/cups.socket'
   '/etc/systemd/system/sockets.target.wants/cups.socket'
ln -s '/lib/systemd/system/cups.path' '
   /etc/systemd/system/multi-user.target.wants/cups.path'
#
# systemctl status cups.service
cups.service - CUPS Printing Service
   Loaded: loaded (/lib/systemd/system/cups.service; enabled)
   Active: inactive (dead) since Tue, 01 May 2015 06:42:38...
   Main PID: 17172 (code=exited, status=0/SUCCESS)
   CGroup: name=systemd:/system/cups.service
```

Notice that the status of `cups.service` changes from `disabled` to `enabled` after using the `enable` option on `systemctl`. Also, notice that the `enable` option simply creates a few symbolic links. You may be tempted to create these links yourself. However, the preferred method is to use the `systemctl` command to accomplish this.

Disabling (removing) a service with systemd

You can use the `disable` option on the `systemctl` command to keep a service from starting at boot. However, it does not immediately stop the service. You need to use the

stop option discussed in the section “Stopping a service with systemd.” The following example shows how to disable a currently enabled service.

```
# systemctl disable cups.service
rm '/etc/systemd/system/printer.target.wants/cups.service'
rm '/etc/systemd/system/sockets.target.wants/cups.socket'
rm '/etc/systemd/system/multi-user.target.wants/cups.path'
#
# systemctl status cups.service
cups.service - CUPS Printing Service
   Loaded: loaded (/lib/systemd/system/cups.service; disabled)
   Active: active (running) since Tue, 01 May 2015 06:06:41...
 Main PID: 17172 (cupsd)
    CGroup: name=systemd:/system/cups.service
            └─ 17172 /usr/sbin/cupsd -f
```

The `disable` option simply removes a few files via the preferred method of the `systemctl` command. Notice also in the preceding example that although the `cups` service is now disabled, the `cups` daemon is still active (running).

The `systemd` `init` has some services that cannot be disabled. These services are static services. Consider the following service, `dbus.service`:

```
# systemctl status dbus.service
dbus.service - D-Bus System Message Bus
   Loaded: loaded (/lib/systemd/system/dbus.service; static)
   Active: active (running) since Mon, 30 Apr 2015 12:35:...
 Main PID: 707 (dbus-daemon)
...
#
# systemctl disable dbus.service
#
# systemctl status dbus.service
dbus.service - D-Bus System Message Bus
   Loaded: loaded (/lib/systemd/system/dbus.service; static)
   Active: active (running) since Mon, 30 Apr 2015 12:35:...
 Main PID: 707 (dbus-daemon)
...

```

When the `systemctl disable` command is issued on `dbus.service`, it is simply ignored. Remember that `static` stands for “statically enabled” and it means that the service is enabled by default and cannot be disabled, even by root. Any services that should not be disabled on your Linux server are set to `static`.

Now that you understand how to enable (or disable) individual services to be persistent, you need to look at service groups as a whole. The next section covers how to start groups of services at boot time.

Configuring a Default runlevel or target unit

Whereas a persistent service is one that is started at server boot time, a persistent (default) `runlevel` or `target unit` is a group of services that are started at boot time. Both classic SysVinit and Upstart define these groups of services as runlevels, while `systemd` calls them `target units`.

Configuring the classic SysVinit daemon default runlevel

You set the persistent runlevel for a Linux server using the SysVinit daemon in the `/etc/inittab` file. A portion of this file is shown here:

```
# cat /etc/inittab
#
# inittab          This file describes how the INIT process should
#                  set up the system in a certain run-level.
...
id:5:initdefault:

# System initialization.
si::sysinit:/etc/rc.d/rc.sysinit
...
```

The bold line in the example shows that the current default runlevel is runlevel 5. To change this, simply edit the `/etc/inittab` file using your favorite editor and change the 5 to one of the following runlevels: 2, 3, or 4. Do *not* use the runlevels 0 or 6 in this file! This would cause your server to either halt or reboot when it is started up.

Configuring the Upstart init daemon default runlevel

Some distributions still use the `/etc/inittab` file to set the default runlevel, whereas others use the `/etc/init/rc-sysinit.conf` file.

Fedora and RHEL's Upstart `init` daemon still uses the `/etc/inittab` file. Therefore, just change the default runlevel as you would on a SysVinit system.

Ubuntu's Upstart `init` daemon uses the `/etc/init/rc-sysinit.conf` file to set the default runlevel, a portion of which is shown in the code that follows. The code line to change is `env DEFAULT_RUNLEVEL=`. Simply edit this file and change that number to the runlevel you desire. However, remember that Ubuntu's runlevel 2 is equivalent to runlevels 3, 4, and 5.

```
$ cat /etc/init/rc-sysinit.conf
# rc-sysinit - System V initialisation compatibility
...
# Default runlevel, this may be overridden on the kernel command-line
# or by faking an old /etc/inittab entry
env DEFAULT_RUNLEVEL=2
```


Configuring the systemd init default target unit

systemd init uses the term “target units” for groups of services to be started. The following shows the various target units you can configure to be persistent and their equivalent backward-compatible, runlevel-specific target units.

- multi-user.target =
 - runlevel2.target
 - runlevel3.target
 - runlevel4.target
- graphical.target = runlevel5.target

The persistent target unit is set via a symbolic link to the default.target unit file. Consider the following:

```
# ls -l /etc/systemd/system/default.target
lrwxrwxrwx. 1 root root 36 Mar 13 17:27
/etc/systemd/system/default.target ->
/lib/systemd/system/runlevel5.target
#
# ls -l /lib/systemd/system/runlevel5.target
lrwxrwxrwx. 1 root root 16 Mar 27 15:39
/lib/systemd/system/runlevel5.target ->
graphical.target
```

The example shows that the current persistent target unit on this server is runlevel5.target because default.target is a symbolic link to the runlevel5.target unit file. However, notice that runlevel5.target is also a symbolic link and it points to graphical.target. Thus, this server’s current persistent target unit is graphical.target.

To set a different target unit to be persistent, you simply need to change the symbolic link for default.target. To be consistent, stick with the runlevel target units if they are used on your server.

The following example changes the server’s persistent target unit from graphical.target to multi-user.target by changing the default.target symbolic link from runlevel5.target to runlevel3.target. The `f` option is used on the `ls -s` command to force any current symbolic link to be broken and the new designated symbolic link to be enforced.

```
# ls -l /lib/systemd/system/runlevel3.target
lrwxrwxrwx. 1 root root 17 Mar 27 15:39
/lib/systemd/system/runlevel3.target ->
multi-user.target
#
# ln -sf /lib/systemd/system/runlevel3.target
/etc/systemd/system/default.target
```

```
#  
# ls -l /etc/systemd/system/default.target  
lrwxrwxrwx. 1 root root 36 May  1 10:06  
  /etc/systemd/system/default.target ->  
  /lib/systemd/system/runlevel3.target
```

When the server is rebooted, the `multi-user.target` will be the persistent target unit. Any services in the `multi-user.target` unit will be started (activated) at that time.

Adding New or Customized Services

Occasionally you will have to add a new service to your Linux server. Also, you may have to customize a particular service. When these needs arise, you must follow specific steps for your Linux server's `init` daemon to either take over the management of the service or recognize the customization of it.

Adding new services to classic SysVinit daemon

When adding a new or customized service to a Linux SysVinit server, you have to complete three steps in order to have the service managed by SysVinit.

1. Create a new or customized service script file.
2. Move the new or customized service script to the proper location for SysVinit management.
3. Add the service to a specific `runlevel(s)`.

Step 1: Create a new or customized service script file

If you are customizing a service script, simply make a copy of the original unit file from `/etc/rc.d/init.d` and add any desired customizations.

If you are creating a new script, you will need to make sure you handle all the various options you want the `service` command to accept for your service, such as `start`, `stop`, `restart`, and so on.

For a new script, especially if you have never created a service script before, it would be wise to make a copy of a current service script from `/etc/rc.d/init.d` and then modify it to meet your new service's needs. Consider the following partial example of the `cupsd` service's script:

```
# cat /etc/rc.d/init.d/cups  
#!/bin/sh  
#  
...
```

```

start () {
    echo -n $"Starting $prog: "

    # start daemon
    daemon $DAEMON
    RETVAL=$?
    echo
    [ $RETVAL = 0 ] && touch /var/lock/subsys/cups
    return $RETVAL
}

stop () {
    # stop daemon
    echo -n $"Stopping $prog: "
    killproc $DAEMON
    RETVAL=$?
    echo
    [ $RETVAL = 0 ] && rm -f /var/lock/subsys/cups
}

restart() {
    stop
    start
}

case $1 in
...

```

The cups daemon's service script starts out by creating functions for each of the start, stop, and restart options. If you feel uncomfortable with shell script writing, review Chapter 7, "Writing Simple Shell Scripts," or consult a resource such as *Linux Command Line and Shell Scripting Bible* (Wiley, 2011) to improve your skills.

Step 2: Move the service script

Once you have modified or created and tested your service's script file, you can move it to the proper location, `/etc/rc.d/init.d`:

```

# cp My_New_Service /etc/rc.d/init.d
#
# ls /etc/rc.d/init.d/My_New_Service
/etc/rc.d/init.d/My_New_Service

```

Step 3: Add the service to runlevels

This final step is needed only if you want the service to be persistent at certain runlevels. You must create a symbolic link for every runlevel at which you want the service to be persistent. Below are the actions needed to accomplish this final step.

1. Check each `runlevel` directory you want the service to start on and determine what the appropriate *S number* should be for your service. For example, the `My_New_Service` should be started at `runlevel 3` and after all the `runlevel 3` services have been started. Therefore, `S100` is an appropriate *S* number for the symbolic filename, as shown here:

```
# ls /etc/rc.d/rc3.d/S*
/etc/rc.d/rc3.d/S00microcode_ctl
/etc/rc.d/rc3.d/S04readahead_early
/etc/rc.d/rc3.d/S05kudzu
...
/etc/rc.d/rc3.d/S98haldaemon
/etc/rc.d/rc3.d/S99firstboot
/etc/rc.d/rc3.d/S99local
/etc/rc.d/rc3.d/S99smartd
# ln -s /etc/rc.d/init.d/My_New_Service
    /etc/rc.d/rc3.d/S100My_New_Service
#
# ls -l /etc/rc.d/rc3.d/S100My_New_Service
lrwxrwxrwx 1 root root 31 2015-05-07 11:01
    /etc/rc.d/rc3.d/S100My_New_Service ->
    /etc/rc.d/init.d/My_New_Service
```

2. Once you have made the symbolic link(s), test that your new or modified service will work as expected before performing a server reboot.

```
# service My_New_Service start
Starting My_New_Service:      [ OK ]
#
# service My_New_Service stop
Stopping My_New_Service:     [ OK ]
```

Once everything is in place, your new or modified service will start at every `runlevel` you have selected on your system. Also, you can start or stop it manually using the `service` command.

Adding new services to the Upstart `init` daemon

You need to complete only one step to add a new service or customize an existing service with the Upstart `init` daemon. Just add a new job configuration file or modify an existing one. However, this one step can be rather complicated.

The Upstart service job configuration files are all located in the `/etc/init` directory. These files are plaintext only. They use a special syntax for directing the Upstart `init` daemon on how to deal with a particular service. The following example of a configuration file has some very simple syntax:

```
# cat ck-log-system-restart.conf
# Upstart event
# ck-log-system-restart - write system restart to log
```

```
#

start on runlevel 6

task
exec /usr/sbin/ck-log-system-restart
```

Any pound sign (#) denotes a comment line and is ignored by Upstart. The other lines are called stanzas and have special syntax for controlling Upstart jobs. The stanzas from the above file are as follows:

- **start on** — This stanza defines what emitted event will start the service or task. In this particular case, when the `runlevel 6` event is emitted, the `ck-log-system-restart` starts.
- **task** — The stanza here defines that this particular job is a task job as opposed to a service.
- **exec** — This stanza defines what program will be run to start the task. Instead of the `exec` stanza, you can embed an actual command line script to run here by using the `script` stanza before the actual code and `end script` after it.

A slightly more complicated job configuration file is shown next — for the `cron` daemon. There are some additional stanzas that were not in the previous example. Notice that the `task` stanza is missing in the file. This indicates that this particular job is a service job instead of a task job.

```
# cat cron.conf
# cron - regular background program processing daemon
#
# cron is a standard UNIX program that runs user-specified
# programs at periodic scheduled times

description    "regular background program processing daemon"

start on runlevel [2345]
stop on runlevel [!2345]

expect fork
respawn

exec cron
```

The additional stanzas in this example are as follows:

- **description** — This stanza is optional and simply describes the service.
- **start on** — Though the `start on` portion of this stanza was previously covered, the `[2345]` syntax was not. Using brackets means that the stanza is valid for any of those numbers. Thus, the service will start on runlevel 2, 3, 4, or 5.

- `stop on` — The stanza here defines what emitted events the service will stop on. The `!2345` in this stanza means *not* `runlevel 2` or `3` or `4` or `5`. In other words, it will only stop on `runlevel 0`, `runlevel 1`, or `runlevel 6`.
- `expect` — This particular stanza is rather important and a little tricky. The `expect fork` syntax will allow Upstart to track this daemon and any of its child processes (forks).
- `respawn` — The stanza here tells Upstart to restart this service should it ever be terminated via a means outside of its normal `stop on`.

Tip

To test your new or modified job configuration files, you can set the `start on` stanza to a non-standard event. In other words, you can make up your own event name. For example, use the event name `MyTest`. To test the new configuration file, you type `initctl emit MyTest` at the command line. If your configuration file works correctly, then modify the `start on` stanza to the correct Upstart event.

Every job configuration file must follow at least three rules. The job configuration file must:

- Not be empty
- Be syntactically correct
- Contain at least one legal stanza

Although there are only three rules, creating or modifying a service job configuration file correctly can be a rather difficult task. See <http://upstart.ubuntu.com/getting-started.html> or <http://upstart.ubuntu.com/cookbook> for help on the syntax needed for these files. Also, you can find out more about events that the Ubuntu `init` daemon emits by typing `man upstart-events` at the command line.

Adding new services to systemd init

When adding a new or customized service to a Linux `systemd` server, you have to complete three steps in order to have the service managed by `systemd`:

1. Create a new or customized service configuration unit file for the new or customized service.
2. Move the new or customized service configuration unit file to the proper location for `systemd` management.
3. Add the service to a specific `target` unit's `Wants` if you want to have the new or customized service start automatically with other services.

Step 1: Create a new or customized service configuration unit file

If you are customizing a service configuration unit file, simply make a copy of the original unit file from `/lib/systemd/system` and add any desired customizations.

For new files, obviously, you will be creating a `service` unit configuration file from scratch. Consider the following basic `service` unit file template. At bare minimum, you need `Description` and `ExecStart` options for a `service` unit configuration file.

```
# cat My_New_Service.service
[Unit]
Description=My New Service

[Service]
ExecStart=/usr/bin/My_New_Service
```

For additional help on customizing or creating a new configuration unit file and the various needed options, you can use the man pages. At the command line, type **man systemd.service** to find out more about the various `service` unit file options.

Step 2: Move the service configuration unit file

Before you move the new or customized service configuration unit file, you need to be aware that there are two potential locations to store service configuration unit files. The one you choose will determine whether or not the customizations take effect and if they remain persistent through software upgrades.

You can place your system service configuration unit file in one of the following two locations:

- `/etc/systemd/system`
 - This location is used to store customized local service configuration unit files.
 - Files in this location are not overwritten by software installations or upgrades.
 - Files here are used by the system *even* if there is a file of the same name in the `/lib/systemd/system` directory.
- `/lib/systemd/system`
 - This location is used to store system service configuration unit files.
 - Files in this location are overwritten by software installations and upgrades.
 - Files here are used by the system *only* if there is *not* a file of the same name in the `/etc/systemd/system` directory.

Thus, the best place to store your new or customized service configuration unit file is in `/etc/systemd/system`.

TIP

When you create a new or customized service, in order for the change to take effect without a server reboot, you will need to issue a special command. At the command line, type **systemctl daemon-reload**.

Step 3: Add the service to the Wants directory

This final step is optional. It needs to be done only if you want your new service to start with a particular `systemd target` unit. For a service to be activated (started) by a particular `target` unit, it must be in that `target` unit's `Wants` directory.

First, add the line **`WantedBy=desired.target`** to the bottom of your service configuration unit file. The following example shows that the desired `target` unit for this new service is `multi-user.target`.

```
# cat /etc/systemd/system/My_New_Service.service
[Unit]
Description=My New Fake Service

[Service]
ExecStart=/usr/bin/My_New_Service

[Install]
WantedBy=multi-user.target
```

To add a new service unit to a `target` unit, you need to create a symbolic link. The following example shows the files located in the `multi-user.target` unit's `Wants` directory. Previously, in the section “Understanding `systemd` init,” the `systemctl` command was used to list out `Wants`, and it is still the preferred method. Notice that in this directory, the files are symbolic links pointing to service unit configuration files in the `/lib/systemd/system` directory.

```
# ls /etc/systemd/system/multi-user.target.wants
abrt-ccpp.service    cups.path            remote-fs.target
abrt-d.service       fcoe.service         rsyslog.service
abrt-oops.service    irqbalance.service  sendmail.service
abrt-vmcore.service  lldpad.service       sm-client.service
atd.service          mcelog.service       sshd-keygen.service
auditd.service       mdmonitor.service    sshd.service
...
#
# ls -l /etc/systemd/system/multi-user.target.wants
total 0
lrwxrwxrwx. 1 root root 37 Nov  2 22:29
  abrt-ccpp.service ->
    /lib/systemd/system/abrt-ccpp.service
lrwxrwxrwx. 1 root root 33 Nov  2 22:29
  abrt-d.service ->
    /lib/systemd/system/abrt-d.service
...
lrwxrwxrwx. 1 root root 32 Apr 26 20:05
  sshd.service ->
    /lib/systemd/system/sshd.service
```


The following illustrates the process of adding a symbolic link file for `My_New_Service`:

```
# ln -s /etc/systemd/system/My_New_Service.service  
/etc/systemd/system/multi-user.target.wants/My_New_Service.service
```

A symbolic link has been created in the `multi-user.target.wants` directory. Now, the new service, `My_New_Service`, will be activated (started) when the `multi-user.target` unit is activated.

Tip

If you want to change the `systemd target` unit for a service, you will need to change the symbol link to point to a new target `Wants` directory location. Use the `ls -sf` command to force any current symbolic link to be broken and the new designated symbolic link to be enforced.

Together, the three steps will get your new or customized service added to a Linux `systemd` server. Remember that at this point, a new service will not be running until a server reboot. To start the new service before a reboot, review the commands in the section “Stopping and Starting Services.”

Summary

How you start and stop services is dependent upon what `init` daemon is used by your Linux server. Before you do any service management, be sure to use the examples in this chapter to help you determine your Linux server’s `init` daemon.

The concepts of starting and stopping services go along with other service management concepts, such as making a service persistent, starting certain services at server boot time, reloading a service, and restarting a service. All of these concepts will be very helpful as you learn about configuring and managing a Linux print server in the next chapter.

Exercises

Refer to the material in this chapter to complete the tasks that follow. If you are stuck, solutions to the tasks are shown in Appendix B (although in Linux, there are often multiple ways to complete a task). Try each of the exercises before referring to the answers. These tasks assume you are running a Fedora or Red Hat Enterprise Linux system (although some tasks will work on other Linux systems as well).

1. Determine which `init` daemon your server is currently using.
2. What `init` daemon is `sshd` using on your Linux server?

3. Determine your server's previous and current `runlevel`.
4. How can you change the default `runlevel` or `target` `unit` on your Linux server?
5. For each `init` daemon, what command(s) will list out services running (or active) on your server?
6. List out the running (or active) services on your Linux server.
7. For each `init` daemon, what command(s) will show a particular service's current status?
8. Show the status of the `cups` daemon on your Linux server.
9. Attempt to restart the `cups` daemon on your Linux server.
10. Attempt to reload the `cups` daemon on your Linux server.

Configuring a Print Server

IN THIS CHAPTER

Understanding printing in Linux

Setting up printers

Using printing commands

Managing document printing

Sharing printers

You can configure your Linux system to use printers that are connected directly to it (via a USB or parallel port) or that are available for printing over the network. Likewise, any printer you configure on your local system can be shared with users on other Linux, Windows, or Mac systems by opening up your printer as a print server.

You configure a printer as a native Linux printer in Fedora, RHEL, Ubuntu, and other Linux systems with the Common UNIX Printing System (CUPS). To configure a printer to work as a Microsoft Windows style of print server, you can use the Samba service in Linux.

This chapter focuses on CUPS. In particular, it shows you the graphical front end to CUPS, called the Printer Configuration window, which comes with Fedora, Red Hat Enterprise Linux, and other Linux distributions. Using Printer Configuration, you can also configure your printers as print servers so people can print to your printer from their own computers.

If you don't have a desktop or want to print from within a shell script, this chapter shows you how to use printing commands. From the command line, print commands such as `lpr` are available for carrying out printing. Commands also exist for querying print queues (`lpq`), manipulating print queues (`lpc`), and removing print queues (`lprm`).

Common UNIX Printing System

CUPS has become the standard for printing from Linux and other UNIX-like operating systems. It was designed to meet today's needs for standardized printer definitions and sharing on Internet Protocol-based networks (as most computer networks are today). Nearly every Linux

distribution today comes with CUPS as its printing service. Here are some of the service's features:

- **IPP** — CUPS is based on the Internet Printing Protocol (<http://www.pwg.org/ipp>), a standard that was created to simplify how printers can be shared over IP networks. In the IPP model, printer servers and clients who want to print can exchange information about the model and features of a printer using HTTP (that is, web content) protocol. A server can also broadcast the availability of a printer so a printing client can easily find a list of locally available printers without configuration.
- **Drivers** — CUPS also standardized how printer drivers are created. The idea was to have a common format that could be used by printer manufacturers so that a driver could work across all different types of UNIX systems. That way, a manufacturer had to create the driver only once to work for Linux, Mac OS X, and a variety of UNIX derivatives.
- **Printer classes** — You can use printer classes to create multiple print server entries that point to the same printer or one print server entry that points to multiple printers. In the first case, multiple entries can each allow different options (such as pointing to a particular paper tray or printing with certain character sizes or margins). In the second case, you can have a pool of printers so that printing is distributed. In this instance, a malfunctioning printer or a printer that is dealing with very large documents won't bring all printing to a halt. CUPS also supports *implicit classes*, which are print classes that form by merging identical network printers automatically.
- **Printer browsing** — With printer browsing, client computers can see any CUPS printers on your local network with browsing enabled. As a result, clients can simply select the printers they want to use from the printer names broadcast on the network, without needing to know in advance what the printers are named and where they are connected. You can turn off the feature to prevent others on the local network from seeing a printer.
- **UNIX print commands** — To integrate into Linux and other UNIX environments, CUPS offers versions of standard commands for printing and managing printers that have been traditionally offered with UNIX systems.

Instead of using the Printer Configuration window, there are other ways you can configure CUPS printing as well:

- **Configuring CUPS from a browser** — The CUPS project itself offers a web-based interface for adding and managing printers. With the `cupsd` service running, type **localhost:631** from a web browser on the computer running the CUPS service to manage printing. (See the section “Using web-based CUPS administration” later in this chapter.)
- **Configuring CUPS manually** — You also can configure CUPS manually (that is, edit the configuration files and start the `cupsd` daemon from the command line). Configuration files for CUPS are contained in the `/etc/cups` directory. In particular, you might be interested in the `cupsd.conf` file, which identifies

permission, authentication, and other information for the printer daemon, and `printers.conf`, which identifies addresses and options for configured printers. Use the `classes.conf` file to define local printer classes.

Coming From Windows

You can print to CUPS from non-UNIX systems as well. For example, you can use a PostScript printer driver to print directly from Windows XP to your CUPS server. You can use CUPS without modification by configuring the XP computer with a PostScript driver that uses `http://printservername:631/printers/targetPrinter` as its printing port.

You may also be able to use the native Windows printer drivers for the printer instead of the PostScript driver. If the native Windows driver does not work out-of-the-box on your CUPS print queue, you can create a Raw Print Queue under CUPS and use that instead. The Raw Print Queue will directly pass through the data from the Windows native print driver to the printer.

To use CUPS, you need to have the `cups` package installed in Fedora or RHEL. Most desktop Linux distributions include CUPS during the initial system install. If it is not installed in a Fedora or RHEL install, install it by typing the following:

```
# yum install cups
```

Setting Up Printers

Although using the printer administration tools specifically built for your distribution is usually best, many Linux systems simply rely on the tools that come with the CUPS software package.

This section explores how to use CUPS web-based administration tools that come with every Linux distribution and then examines the printer configuration tool `system-config-printer`, which comes with Fedora and Red Hat Enterprise Linux systems to enable you to set up printers. In some cases, no configuration is necessary, because connected printers can be automatically detected and configured.

Adding a printer automatically

CUPS printers can be configured to automatically broadcast their availability on the network so a client system can detect and use them without configuration. Connect a USB printer to your computer, and the printer can be automatically detected and made available. In fact, if you attach a local printer in Fedora and the print driver is not yet installed, you are prompted to install the software packages needed to use the printer.

The first time you go to print a document or view your printer configuration tool, the printers will be there ready to use. Further configuration can be done using the web-based CUPS administration tool or the Printer Configuration window.

Using web-based CUPS administration


CUPS offers its own web-based administrative tool for adding, deleting, and modifying printer configurations on your computer. The CUPS print service (using the `cupsd` daemon) listens on port 631 to provide access to the CUPS web-based administrative interface and share printers.

If CUPS is already running on your computer, you can immediately use CUPS web-based administration from your web browser. To see whether CUPS is running and to start setting up your printers, open a web browser on the local computer and type the following into its location box: **`http://localhost:631/`**.

A prompt for a valid login name and password may appear when you request functions that require it. If so, type the root login name and the root user's password, and then click OK. A screen similar to the one shown in Figure 16.1 appears.

FIGURE 16.1

CUPS provides a web-based administration tool.



By default, web-based CUPS administration is available only from the local host. To access web-based CUPS administration from another computer, from the main CUPS page, select the Administration tab, select the check box next to Allow remote administration, and select the Change Settings button. Then, from a remote browser, you can access the CUPS Administration page by going to port 631 on the CUPS server (for example, `http://host.example.com:631`).

To configure a printer that is not automatically detected, you can add a printer from the Admin screen. With the Admin screen displayed, you can add a printer as follows:

1. **Click the Add Printer button.** The Add New Printer screen appears.
2. **Select the device to which the printer is connected.** The printer can be connected locally to a parallel, SCSI, serial, or USB port directly on the computer. Alternatively, you can select a network connection type for Apple printers (appSocket/HP JetDirect), Internet Printing Protocol (`http` or `ipp`), or a Windows printer (using Samba or SMB).
3. **If prompted for more information, you may need to further describe the connection to the printer.** For example, you may need to enter the baud rate and parity for a serial port, or you might be asked for the network address for an IPP or Samba printer.
4. **Type a Name, Location, and Description for the printer and select if you want to share this printer.** Then click Continue.
5. **Select the make of the print driver.** If you don't see the manufacturer of your printer listed, choose PostScript for a PostScript printer or HP for a PCL printer. For the manufacturer you choose, you will be able to select a specific model.
6. **Set options.** If you are asked to set options for your printer, you may do so. Then select Set Printer Options to continue.
7. **Your printer should be available.** If the printer is added successfully, click the name of your printer to have the new printer page appear; from the printer page, you can select Maintenance or Administration to print a test page or modify the printer configuration.

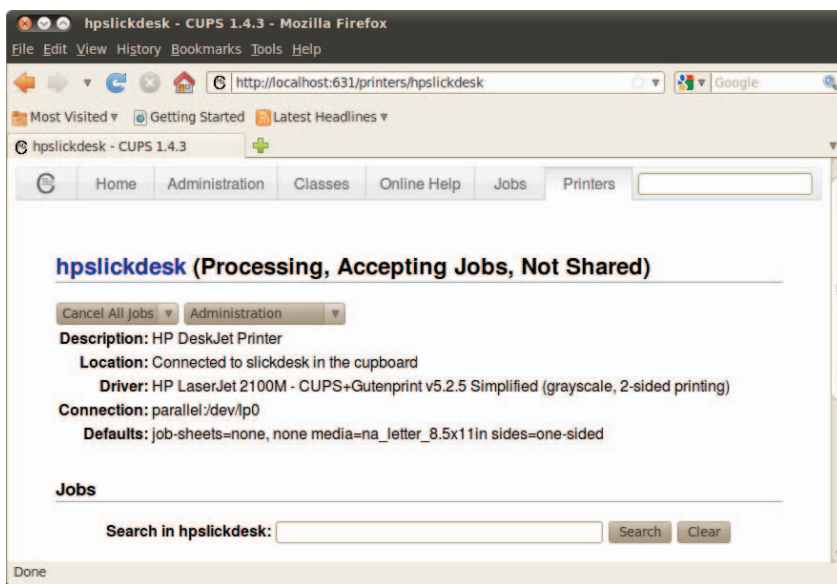
With the basic printer configuration done, you can now do further work with your printers. Here are a few examples of what you can do:

- **List print jobs.** Click Show All Jobs to see what print jobs are currently active from any of the printers configured for this server. Click Show Completed Jobs to see information about jobs that are already printed.
- **Create a printer class.** Click the Administration tab and choose Add Class and identify a name, description, and location for a printer class. From the list of Printers (Members) configured on your server, select the ones to go into this class.

- **Cancel a print job.** If you print a 100-page job by mistake or if the printer is spewing out junk, the Cancel feature can be very handy. From the Administration tab, click Manage Jobs; then click Show Active Jobs to see what print jobs are currently in the queue for the printer. Select the Cancel Job button next to the print job you want to cancel.
- **View printers.** You can click the Printers tab from the top of any of the CUPS web-based administration pages to view the printers you have configured. For each printer that appears, you can select Maintenance or Administrative tasks. Under Maintenance, click Pause Printer (to stop the printer from printing but still accept print jobs for the queue), Reject Jobs (to not accept any further print jobs for the moment), or Print Test Page (to print a page). Figure 16.2 shows the information on the Printers tab.

FIGURE 16.2

You can do administration tasks from the Printers tab.



Using the Printer Configuration window

If you are using Fedora, RHEL, or other Red Hat–based systems, you can use the printer configuration window to set up your printers. In fact, I recommend that you use it instead of CUPS web administration because the resulting printer configuration files are tailored to work with the way the CUPS service is started on those systems.

To install a printer from your GNOME desktop, start the Printer Configuration window by selecting System ⇨ Administration ⇨ Printing, or as root user by typing

system-config-printer. This tool enables you to add and delete printers and edit printer properties. It also enables you to send test pages to those printers to make sure they are working properly.

The key here is that you are configuring printers that are managed by your print daemon (`cupsd` for the CUPS service). After a printer is configured, users on your local system can use it. You can refer to the section “Configuring Print Servers” to learn how to make the server available to users from other computers on your network.

The printers that you set up can be connected directly to your computer (as on a parallel port) or to another computer on the network (for example, from another UNIX system or Windows system).

Configuring local printers with the Printer Configuration window

Add a local printer (in other words, a printer connected directly to your computer) with the Printer Configuration window using the procedure that follows.

Adding a local printer

To add a local printer from a GNOME desktop in Fedora 16, follow these steps:

1. Select Applications ⇨ Other ⇨ Printing or type the following as root user from a Terminal window:

```
# system-config-printer &
```

The Printing window appears.
2. Click Add. (If asked, select the button to Adjust Firewall to allow access to the printer port 631.) A New Printer window appears.
3. If the printer you want to configure is detected, simply select it and click Forward. If it is not detected, choose the device to which the printer is connected (LPT #1 and Serial Port #1 are the first parallel and serial ports, respectively) and click Forward. (Type `/usr/sbin/lpinfo -v | less` in a shell to see printer connection types.) You are asked to identify the printer's driver.
4. To use an installed driver for your printer, choose Select Printer From Database, and then choose the manufacturer of your printer. As an alternative, you could select Provide PPD File and supply your own PPD file (for example, if you have a printer that is not supported in Linux and you have a driver that was supplied with the printer). PPD stands for PostScript Printer Description. Select Forward to see a list of printer models from which you can choose.


Tip

If your printer doesn't appear on the list but supports PCL (HP's Printer Control Language), try selecting one of the HP printers (such as HP LaserJet). If your printer supports PostScript, select PostScript printer from the list. Selecting Raw Print Queue enables you to send documents that are already formatted for a particular printer type to a specific printer.

5. With your printer model selected, click the driver you want to use with that printer. Click Forward to continue.
6. Add the following information and click Forward:
 - **Printer Name** — Add the name you want to give to identify the printer. The name must begin with a letter, but after the initial letter, it can contain a combination of letters, numbers, dashes (-), and underscores (_). For example, an HP printer on a computer named maple could be named hp-maple.
 - **Description** — Add a few words describing the printer, such as its features (for example, an HP LaserJet 2100M with PCL and PS support).
 - **Location** — Add some words that describe the printer's location (for example, "In Room 205 under the coffeepot").
7. When the printer is added, you may be prompted to print a test page (click No or Yes). The new printer entry appears in the Printing window, as shown in Figure 16-3.

FIGURE 16.3

The Printer Properties window after adding a printer



8. If you want the printer to be your default printer, right-click the printer and select Set As Default. As you add other printers, you can change the default printer by selecting the one you want and Set As Default again.
9. Printing should be working at this point. To make sure, open a Terminal window and use the `lpr` command to print a file (such as `lpr/etc/hosts`). (If you want to share this printer with other computers on your network, refer to the section "Configuring Print Servers," later in this chapter.)

Editing a local printer

After double-clicking the printer you want to configure, choose from the following menu options to change its configuration:

- **Settings** — The Description, Location, Device URI, and Make and Model information you created earlier are displayed in this dialog box.
- **Policies** — Click Policies to set the following items:
 - **State** — Select check boxes to indicate whether the printer will print jobs that are in the queue (Enabled), accept new jobs for printing (Accepting Jobs), or be available to be shared with other computers that can communicate with your computer (Shared). You also must select Server Settings and click the Share Published printers connected to this system check box before the printer will accept print jobs from other computers.
 - **Policies** — In case of error, the stop-printer selection causes all printing to that printer to stop. You can also select to have the job discarded (abort-job) or retried (retry-job) in the event of an error condition.
 - **Banner** — There are no starting or ending banner pages by default for the printer. Choose starting or ending banner pages that include text such as Classified, Confidential, Secret, and so on.
- **Access Control** — If your printer is a shared printer, you can select this window to create a list that either allows users access to the printer (with all others denied) or denies users access to the printer (with all others allowed).
- **Printer Options** — Click Printer Options to set defaults for options related to the printer driver. The available options are different for different printers. Many of these options can be overridden when someone prints a document. Here are a few of the options you might (or might not) have available:
 - **Watermark** — Several Watermark settings are available to enable you to add and change watermarks on your printed pages. By default, Watermark and Overlay are off (None). By selecting Watermark (behind the text) or Overlay (over the text), you can set the other Watermark settings to determine how watermarks and overlays are done. Watermarks can go on every page (All) or only the first page (First Only).
 Select Watermark Text to choose what words are used for the watermark or overlay (Draft, Copy, Confidential, Final, and so on). You can then select the font type, size, style, and intensity of the watermark or overlay.
 - **Resolution Enhancement** — You can use the printer's current settings or choose to turn resolution enhancement on or off.
 - **Page Size** — The default is U.S. letter size, but you can also ask the printer to print legal size, envelopes, ISO A4 standard, or several other page sizes.
 - **Media Source** — Choose which tray to print from. Select Tray 1 to insert pages manually.